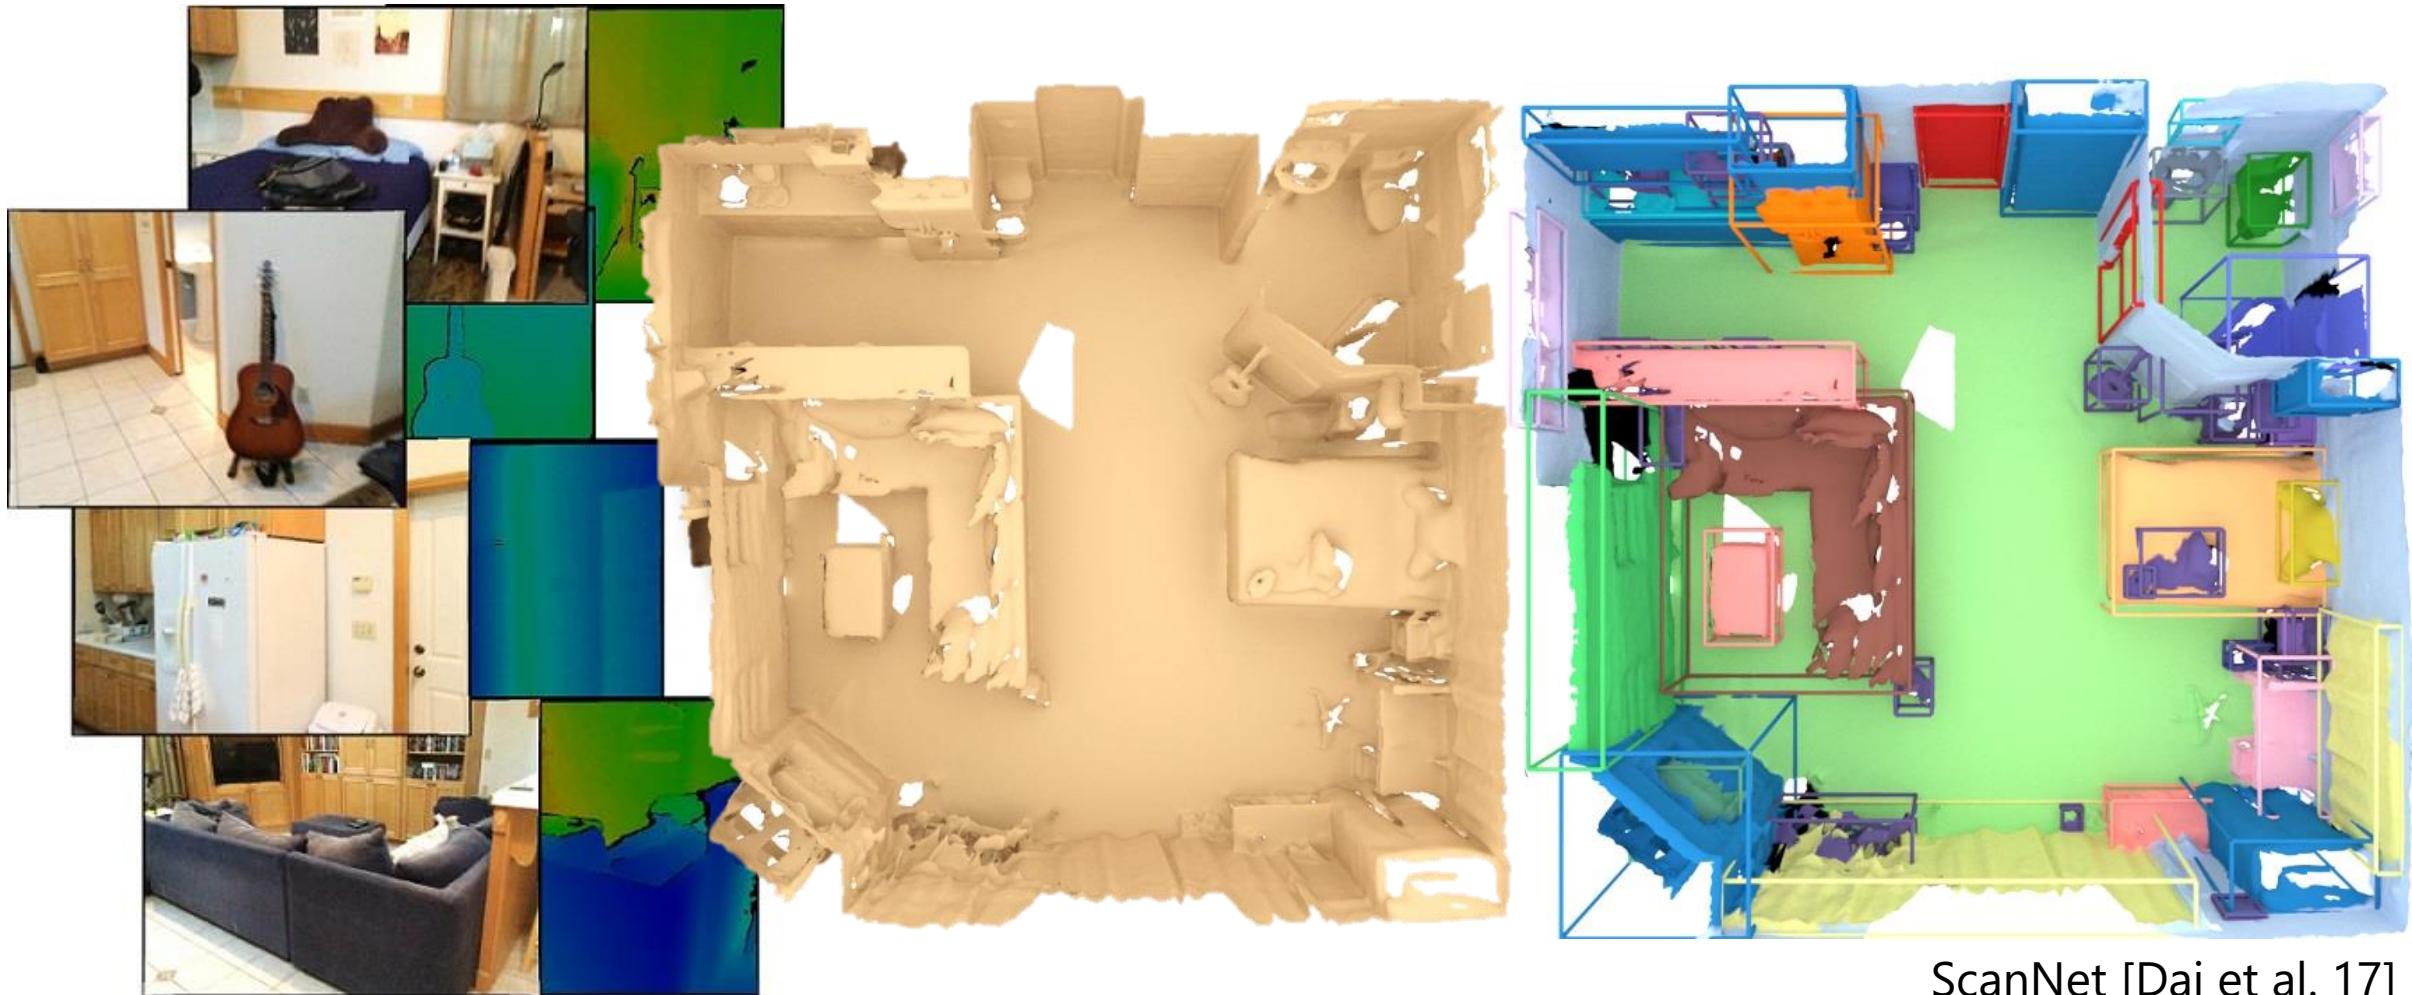


Geometric Foundations: Surface Representations

Prof. Angela Dai

Brief Recap

Machine Perception of Real-World Environments



ScanNet [Dai et al. 17]

We perceive and interact with a 3D world



ASIMO, Honda



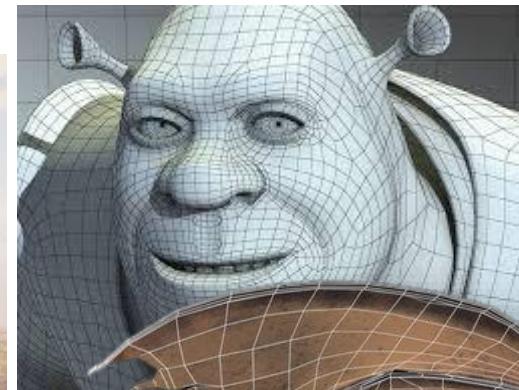
Star Trek TNG (Phantasms)

How to represent 3D?

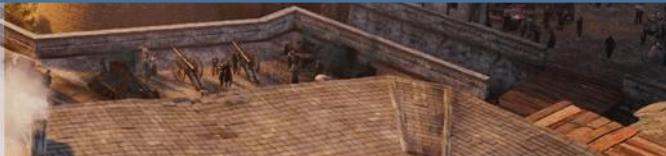
Looking at source and target applications

Looking at source and target applications

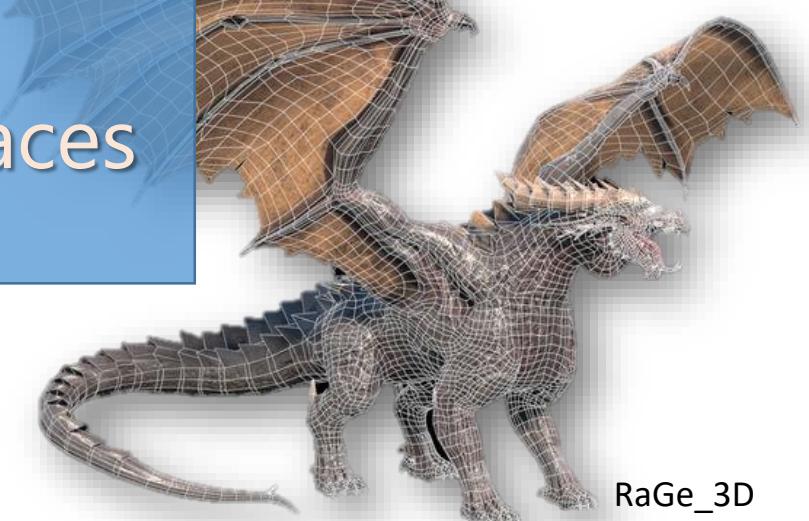
- Movies and games



Need to be able to control and modify surfaces



Assassin's Creed



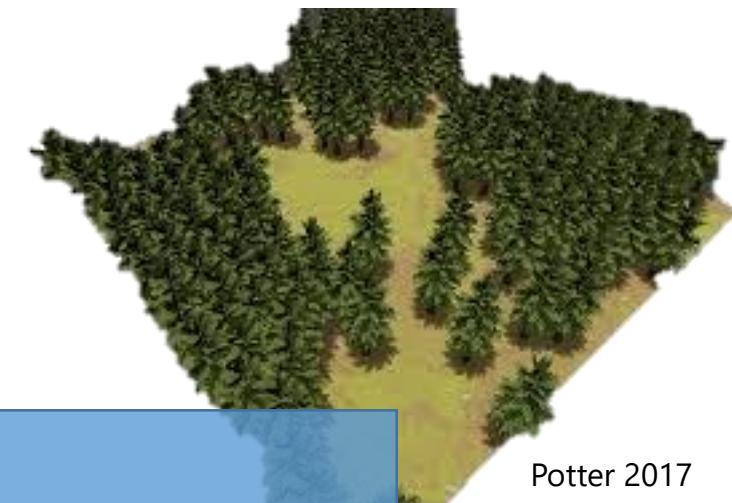
RaGe_3D

Looking at source and target applications

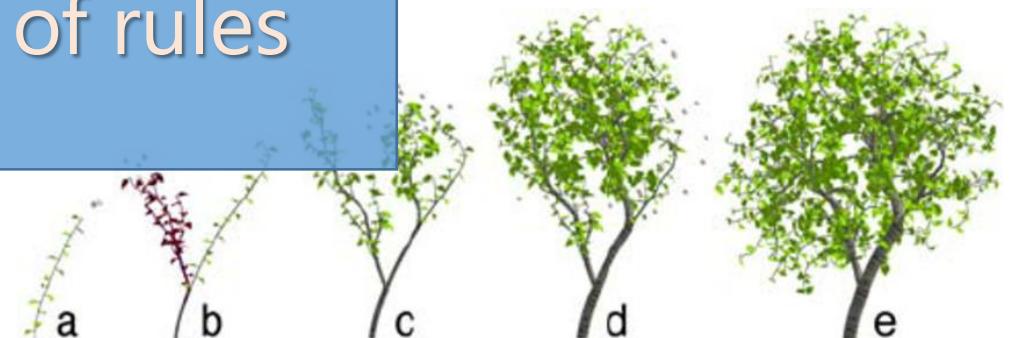
- Movies and games: backgrounds



Generate large scenes from small set of rules



Potter 2017



[Palubicki et al. '09]

<https://www.pluralsight.com/courses/houdini-python-procedural-cities>

Looking at source and target applications

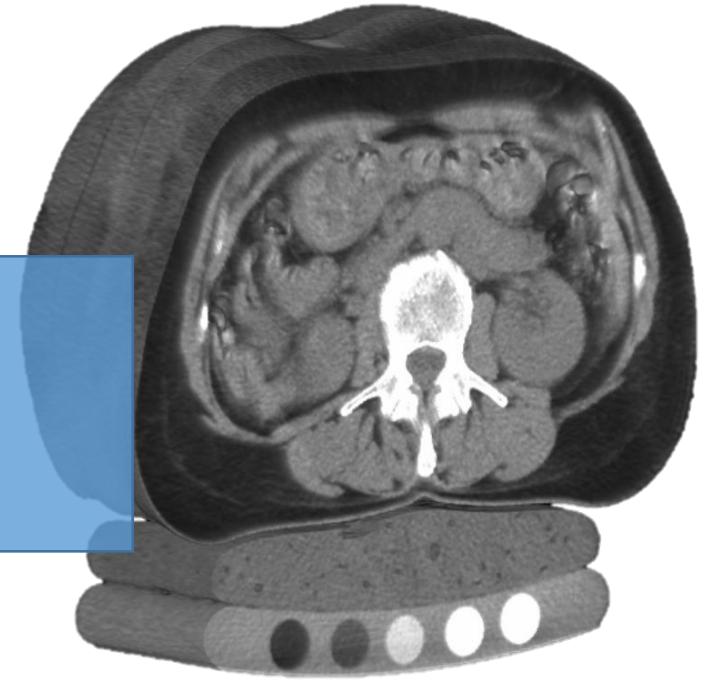
- Real-world data capture



Digital Michelangelo



Waymo Open Dataset

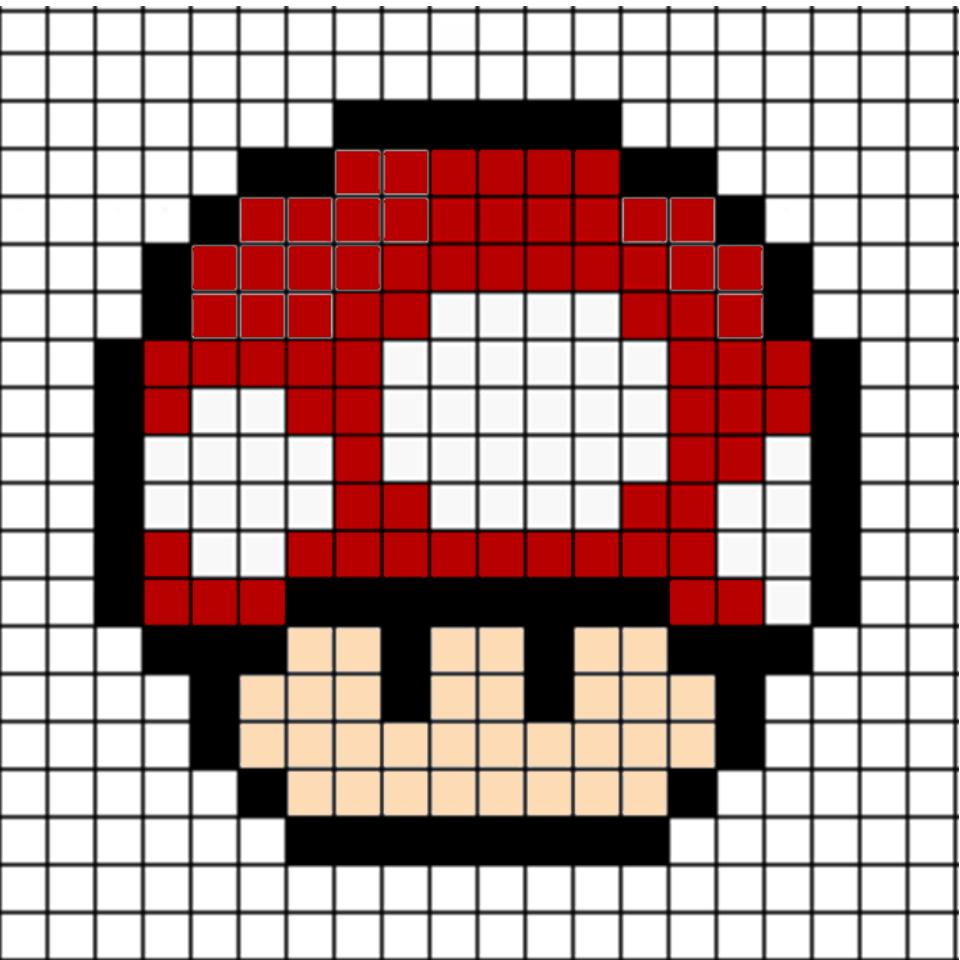
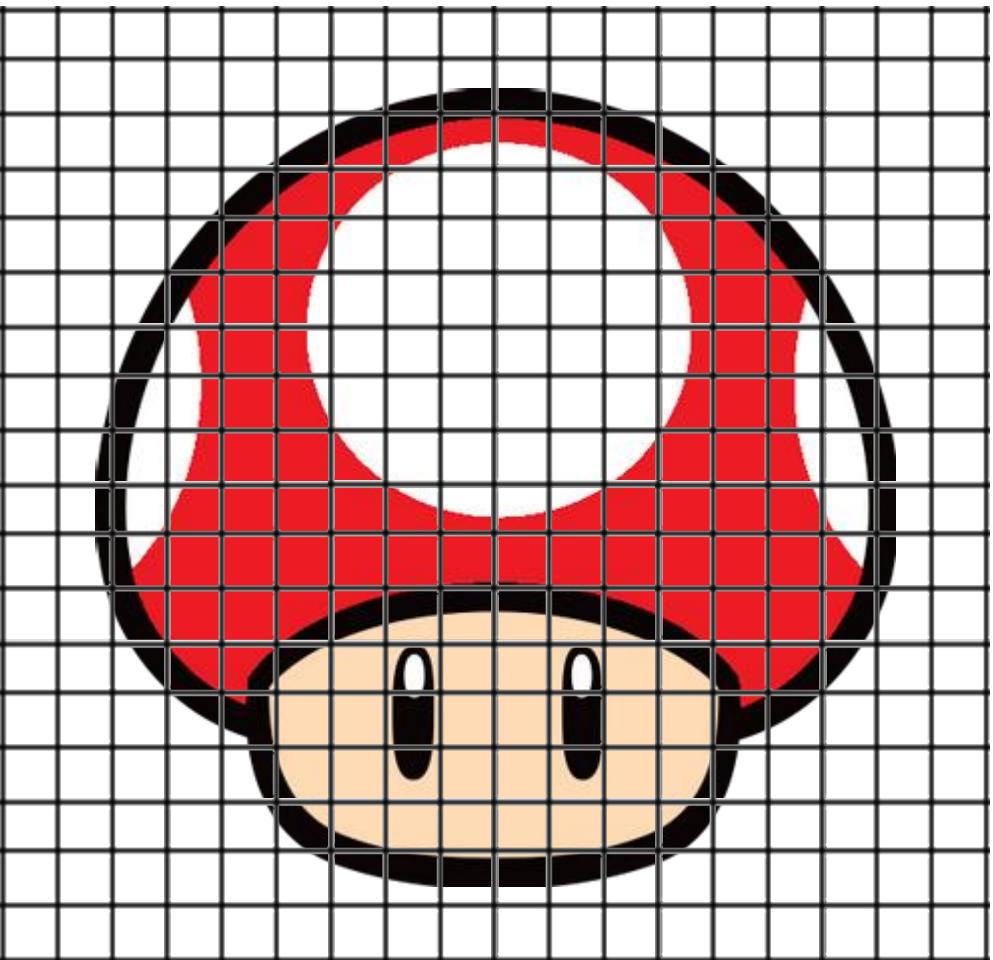


MindwaysCT Software

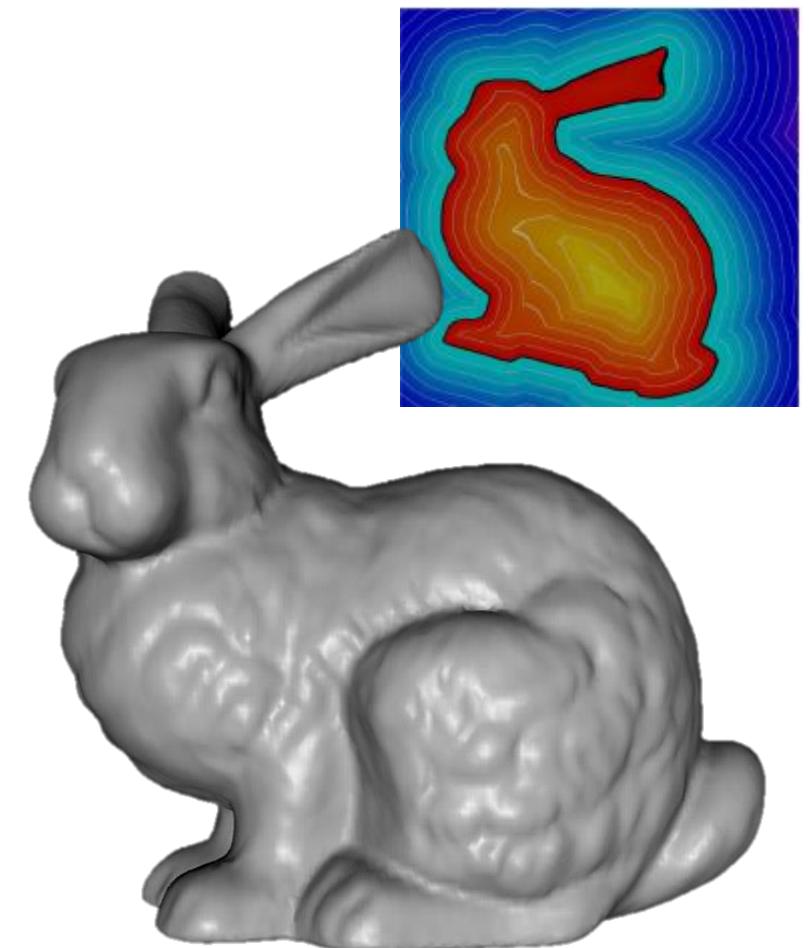
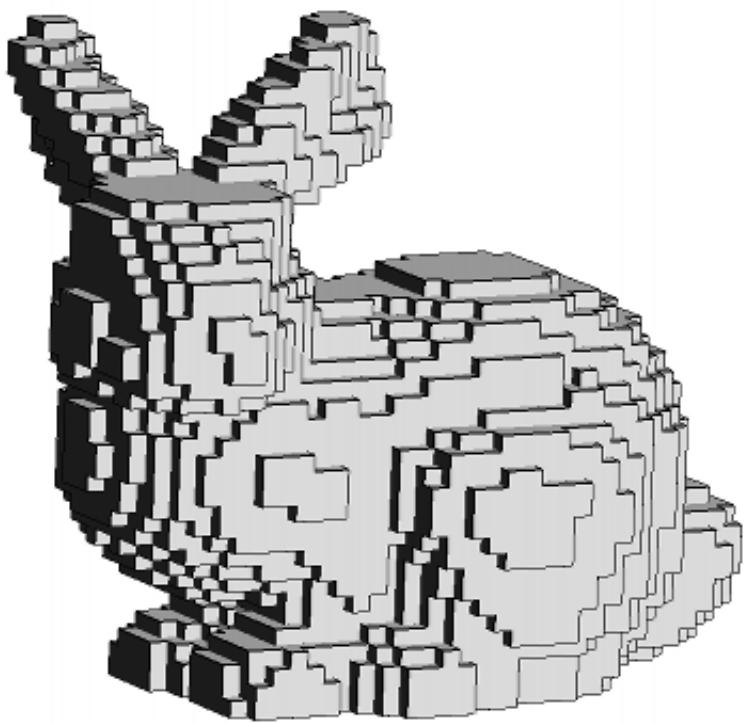
Various 3D Representations

- Things to consider:
- How well does it fit with the constraints of the source data or target application?
- How much memory is necessary for storage?
- Efficiency of operations (editing, transforms, rendering, etc.)

2D Image Representation



3D Voxels



3D Voxels

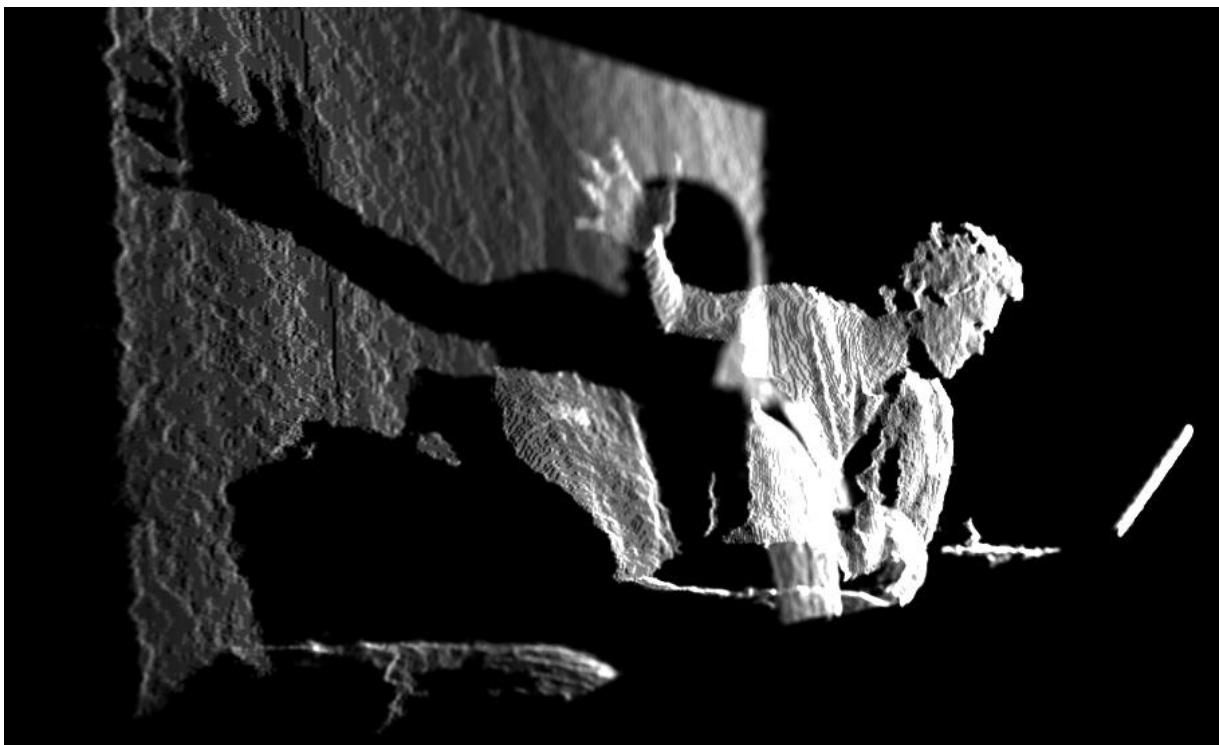
- Regular grid structure (convolutions)
- Store attributes in each grid voxel
 - occupancy, distance, color, normals, etc.
- Represent arbitrary topologies
- Easy to query, operate on neighbors
- Cubic growth for high resolutions



sketchfab - elbriga

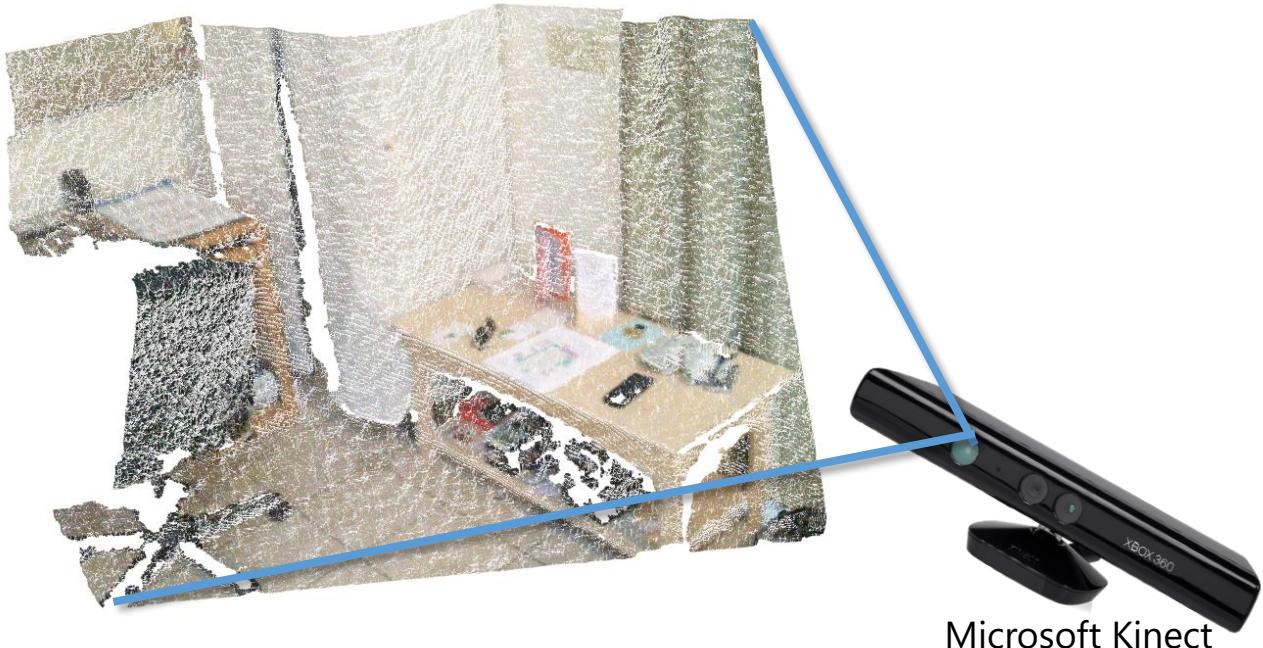
Point Clouds

- Set of points as (x,y,z) locations
- (optionally other attributes, e.g., color)



Point Clouds

- Unordered set of points as (x,y,z) locations
- (optionally other attributes, e.g., color)
- Can represent raw data capture



Microsoft Kinect



Digital Michelangelo Project

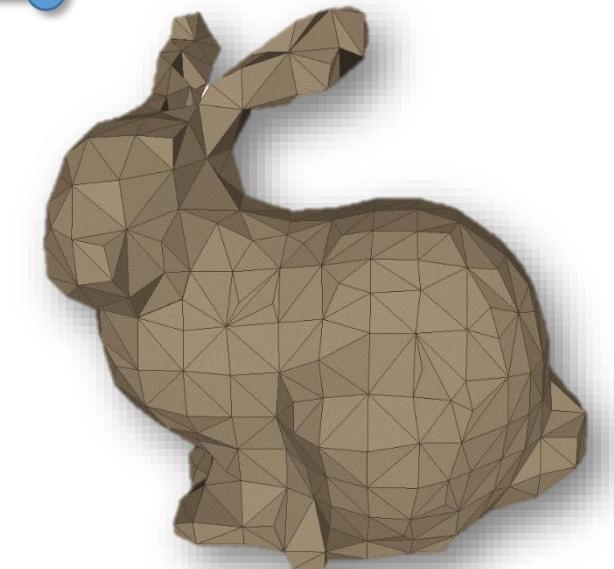
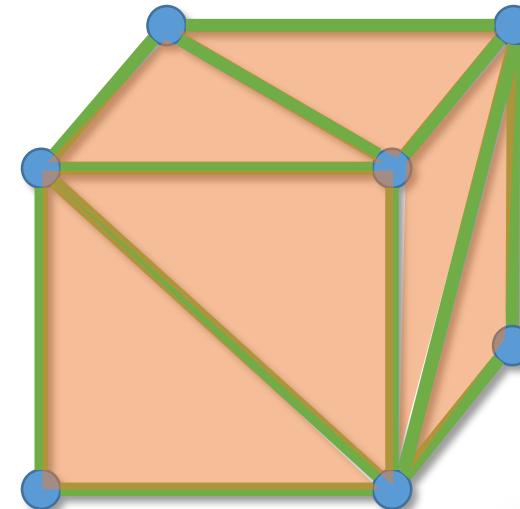
Point Clouds

- Unordered set of points as (x,y,z) locations
- (optionally other attributes, e.g., color)
- Can represent raw data capture
- Can more efficiently represent surface than dense grids
- No spatial structure
- Not so efficient neighbor queries



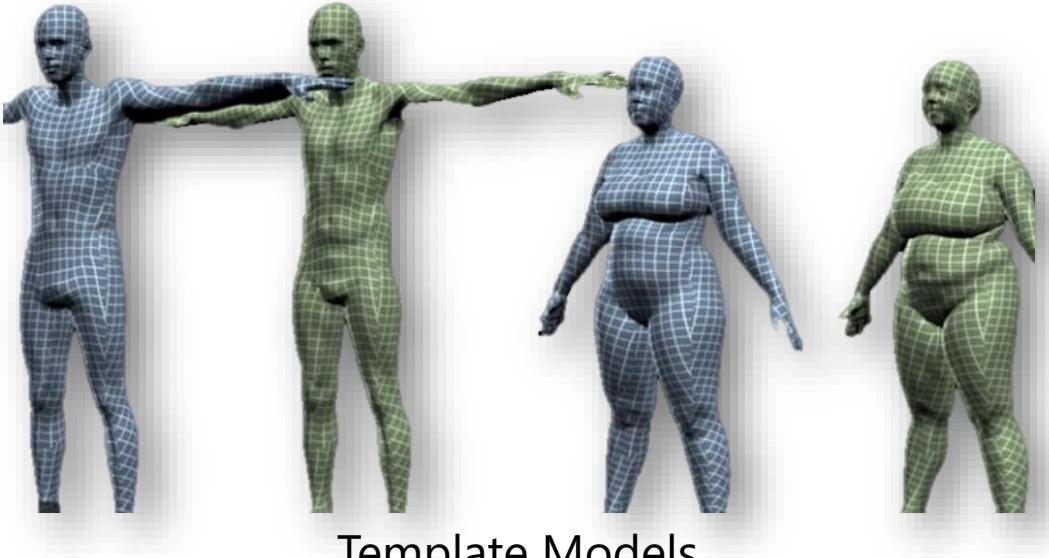
Polygon Meshes

- Collection of **vertices**, **edges**, **faces**
 - Often triangles
- Vertex: 3D position (x, y, z)
- Edge: connection between two vertices
- Face: closed set of edges

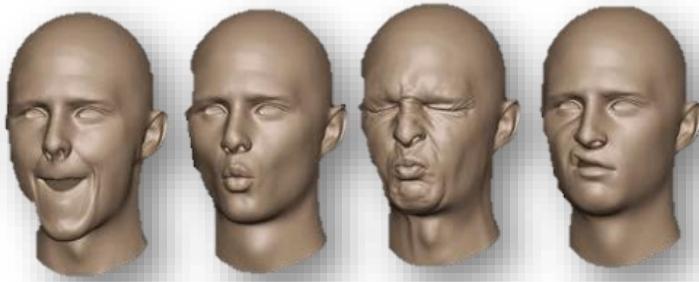


Polygon Meshes

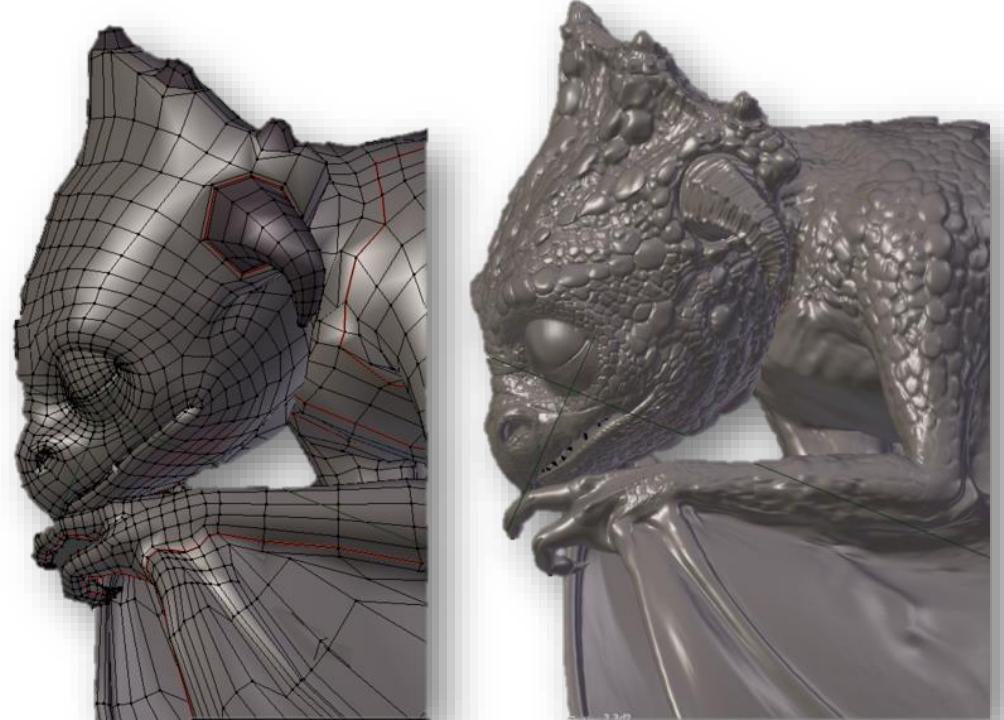
- Many applications



Template Models



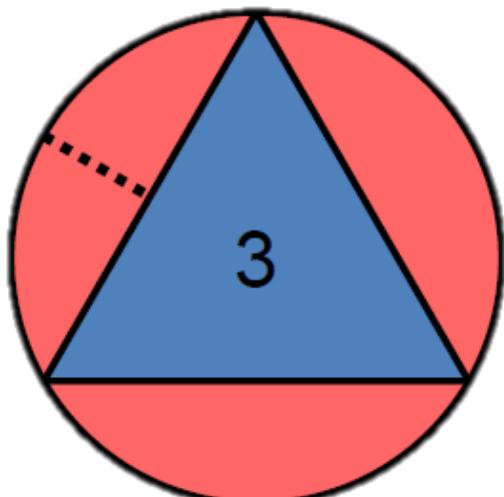
Content Creation: Movies, Games, Mixed Reality



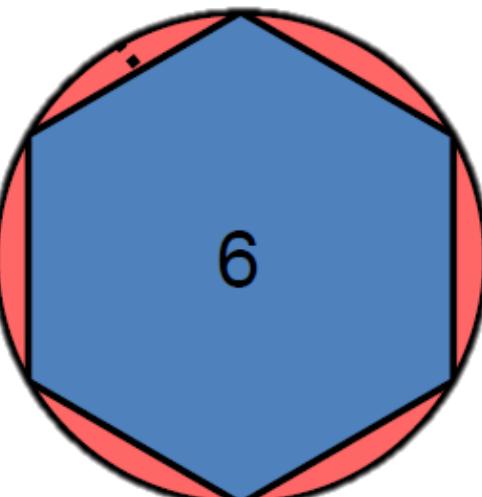
Sintel

Polygon Meshes

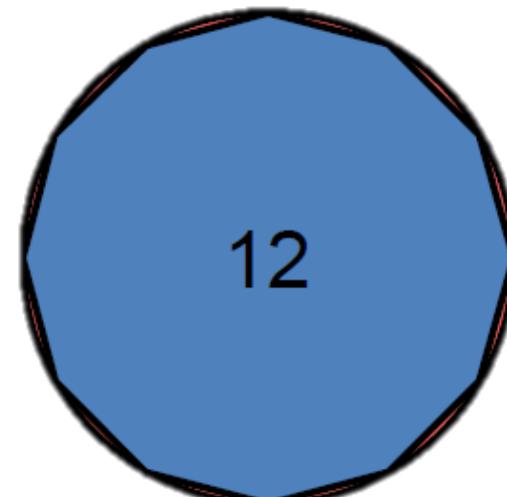
- Piecewise linear approximation of a surface
 - $O(h^2)$ error



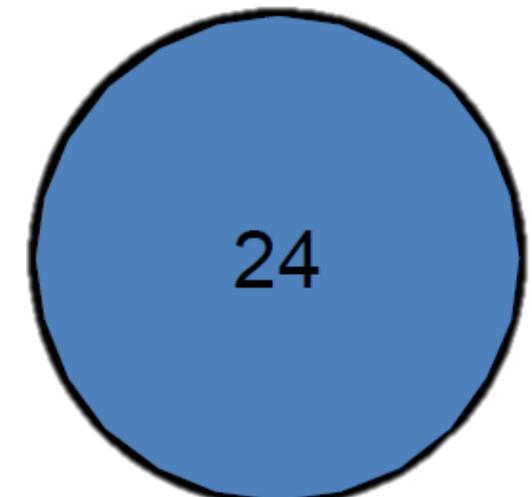
25%



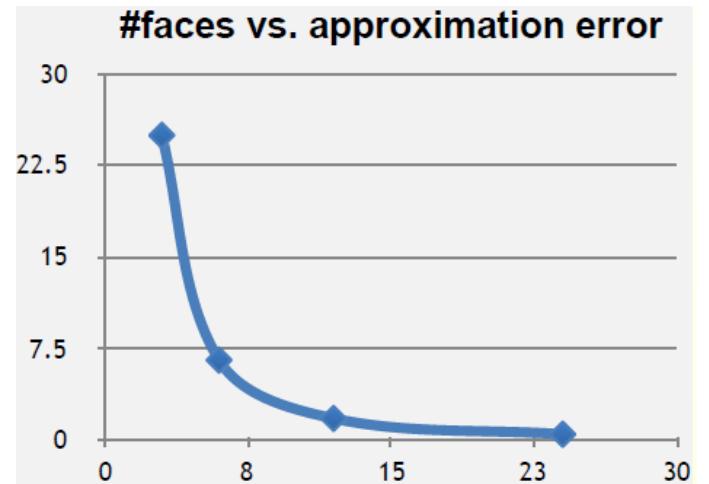
6.5%



1.7%

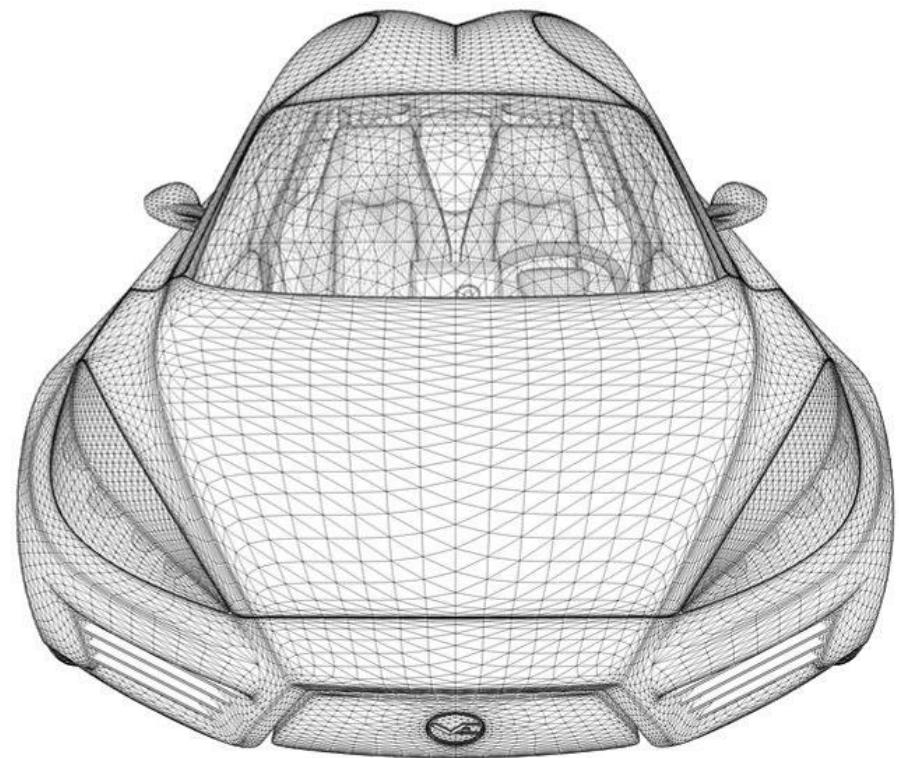


0.4%



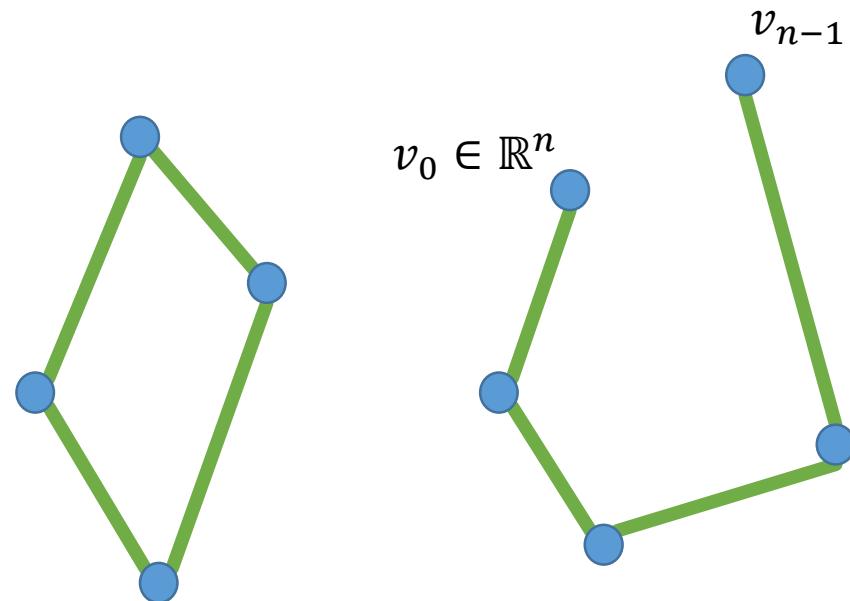
Polygon Meshes

- Popular representation
- Can represent arbitrary topologies
- Can be edited/manipulated
- Efficient rendering



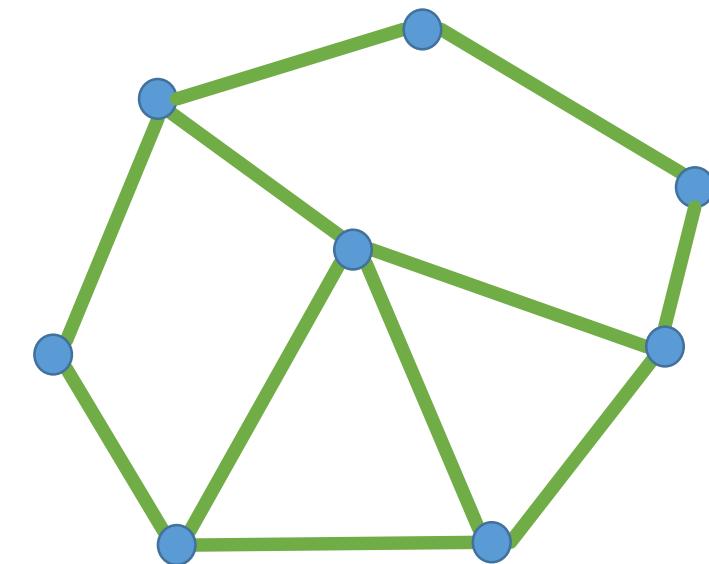
Polygon

- Vertices: v_0, v_1, \dots, v_{n-1}
- Edges: $\{(v_0, v_1), \dots, (v_{n-2}, v_{n-1})\}$
- Closed: $v_0 = v_{n-1}$
- Planar: all vertices on a plane
- Simple: not self-intersecting



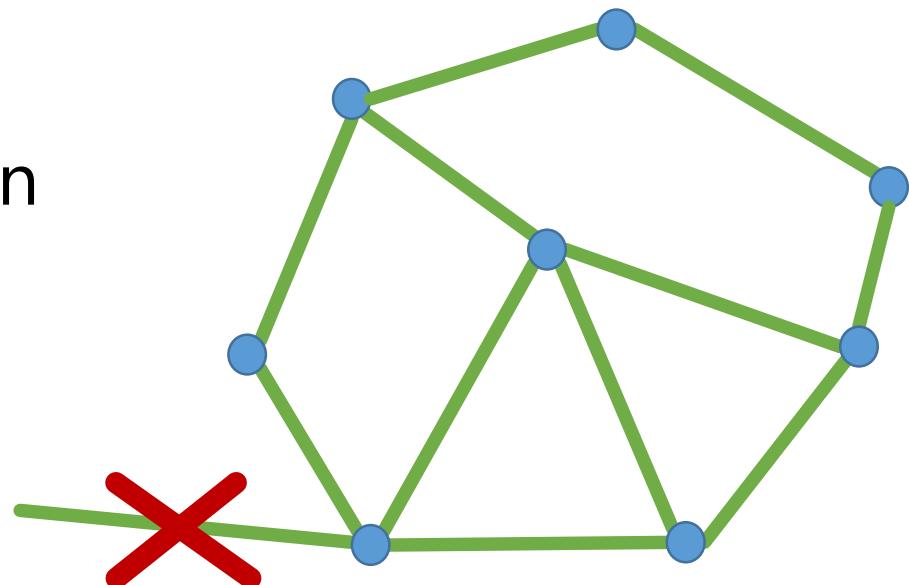
Polygon Mesh

- Polygon Mesh: finite set M of closed, simple polygons P_i
$$M = \langle V, E, F \rangle$$
- Intersection of two polygons in M is either empty, a vertex, or an edge



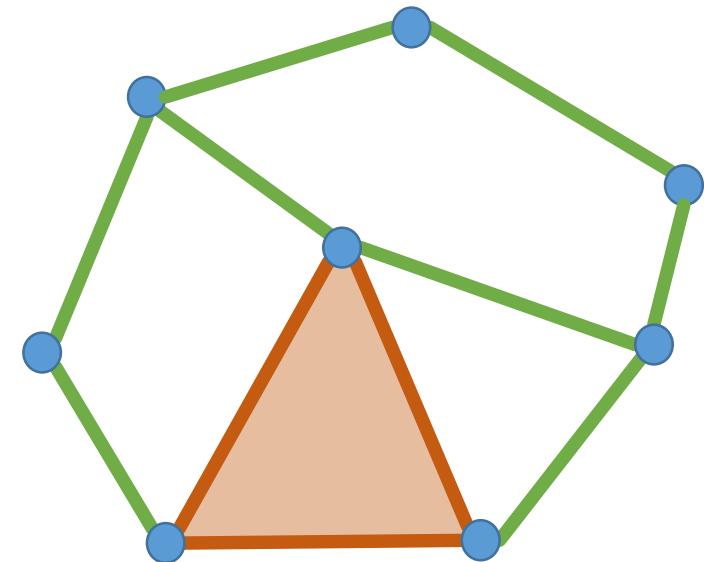
Polygon Mesh

- Polygon Mesh: finite set M of closed, simple polygons P_i
$$M = \langle V, E, F \rangle$$
- Intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon



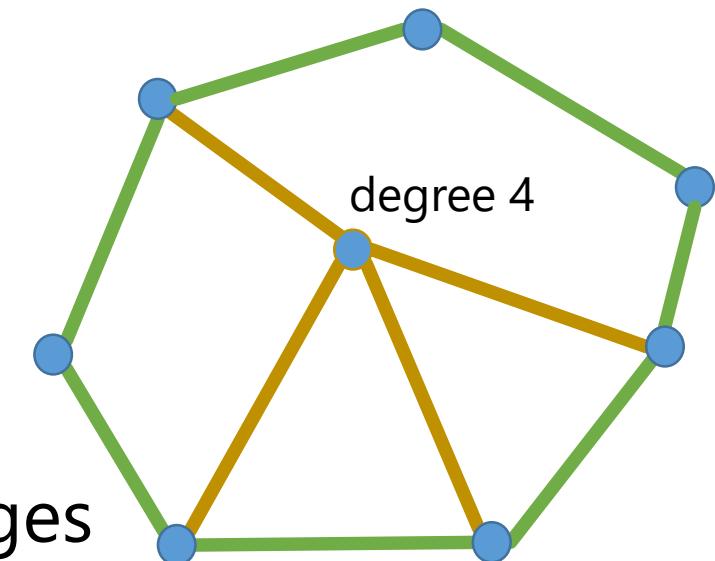
Polygon Mesh

- Polygon Mesh: finite set M of closed, simple polygons P_i
$$M = \langle V, E, F \rangle$$
- Intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each P_i defines a *face* of the polygonal mesh



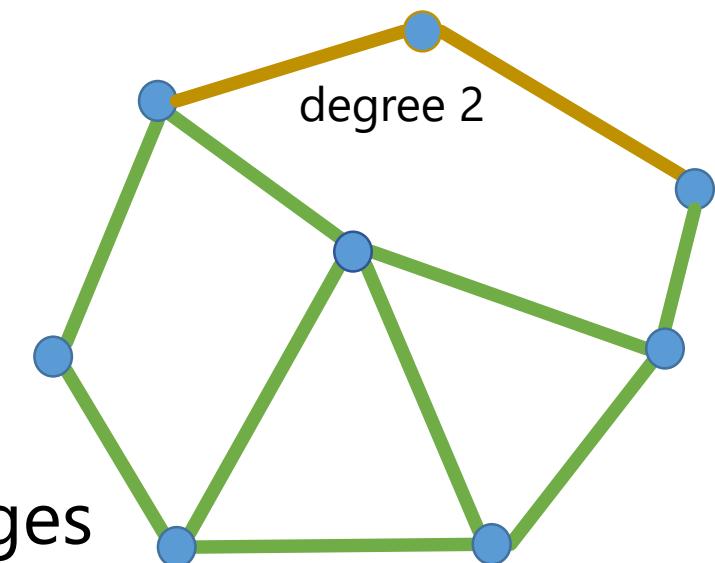
Polygon Mesh

- Polygon Mesh: finite set M of closed, simple polygons P_i
$$M = \langle V, E, F \rangle$$
- Intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each P_i defines a *face* of the polygonal mesh
- Vertex *degree* or *valence*: number of incident edges



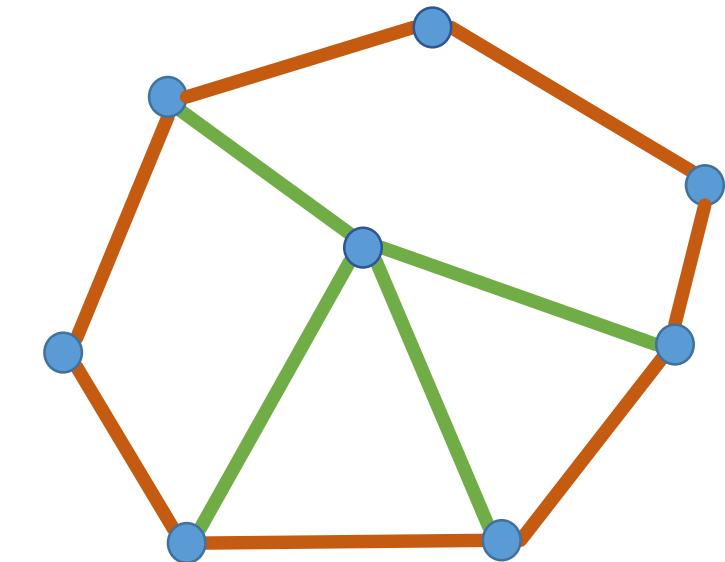
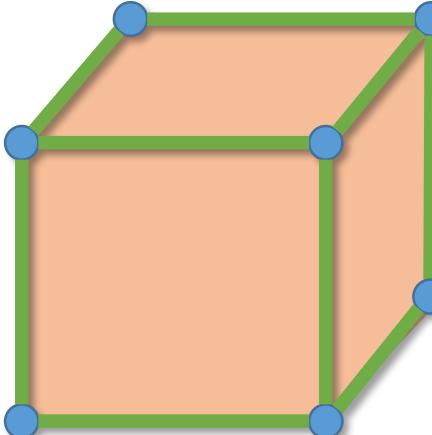
Polygon Mesh

- Polygon Mesh: finite set M of closed, simple polygons P_i
$$M = \langle V, E, F \rangle$$
- Intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each P_i defines a *face* of the polygonal mesh
- Vertex *degree* or *valence*: number of incident edges



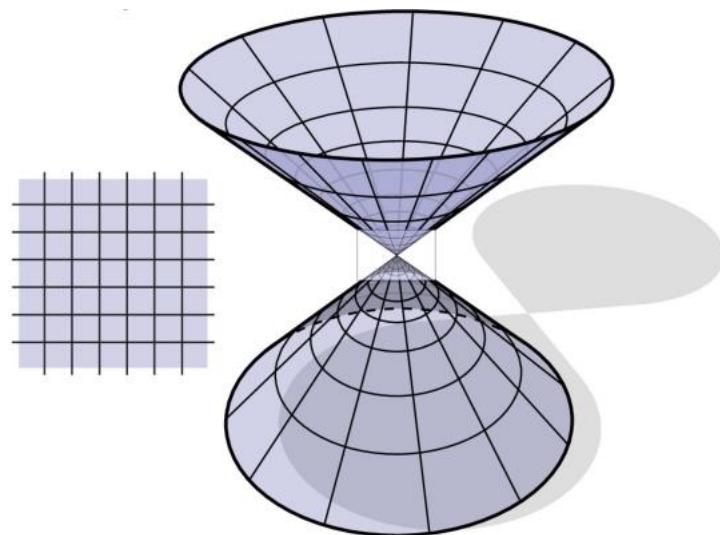
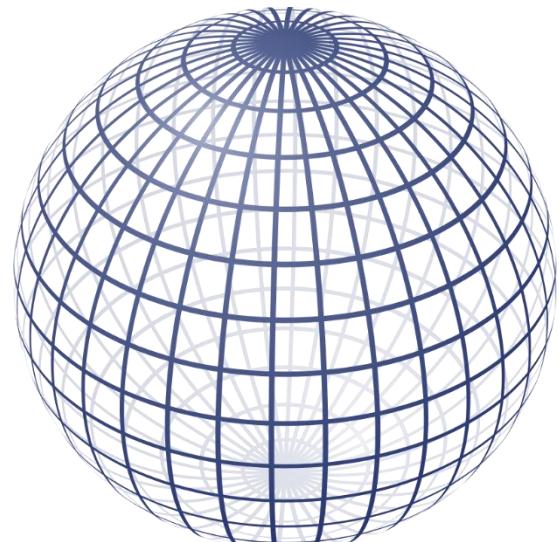
Polygon Mesh

- Polygon Mesh: finite set M of closed, simple polygons P_i
$$M = \langle V, E, F \rangle$$
- Boundary: set of all edges that belong only to one polygon
 - Either empty or forms closed loops
 - If empty, then the polygon mesh is closed

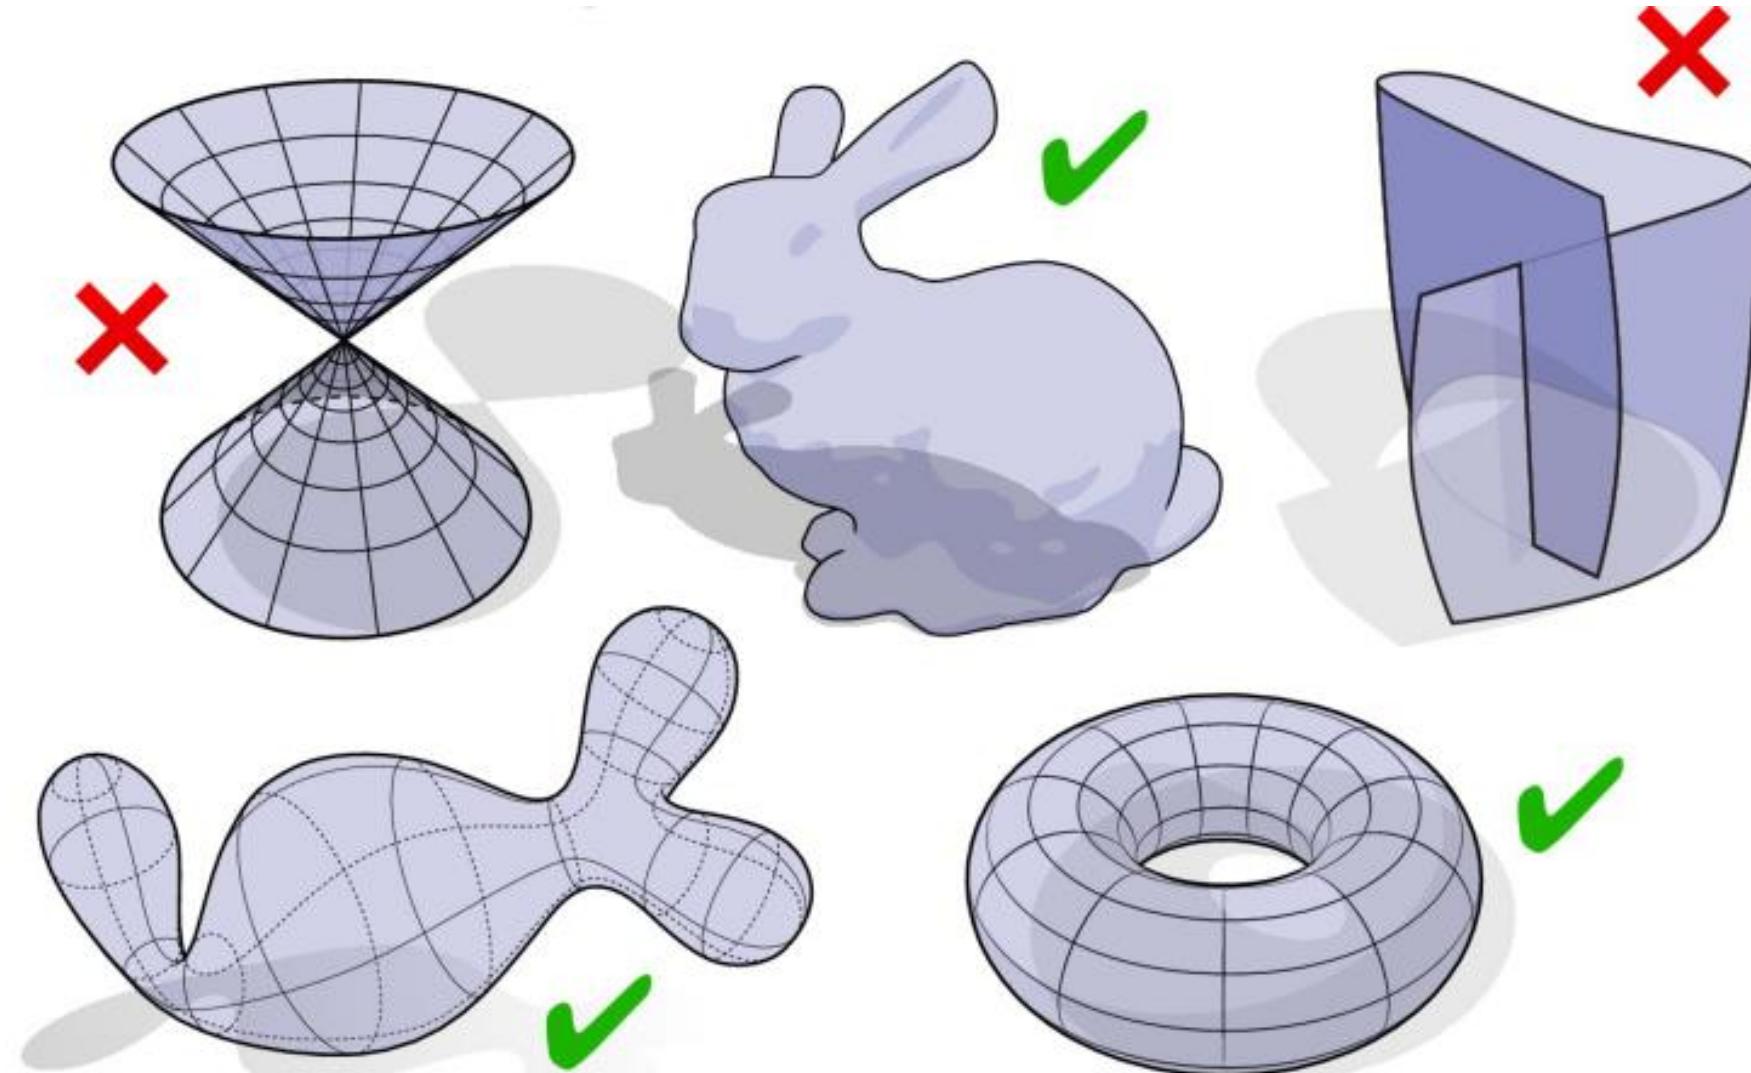


Manifold Geometry

- Recall: we aim to represent surfaces, e.g., “boundary” of an object
- Surfaces are manifold: locally Euclidean
 - i.e., around every point, there is a neighborhood topologically the same as the open unit ball in \mathbb{R}^n

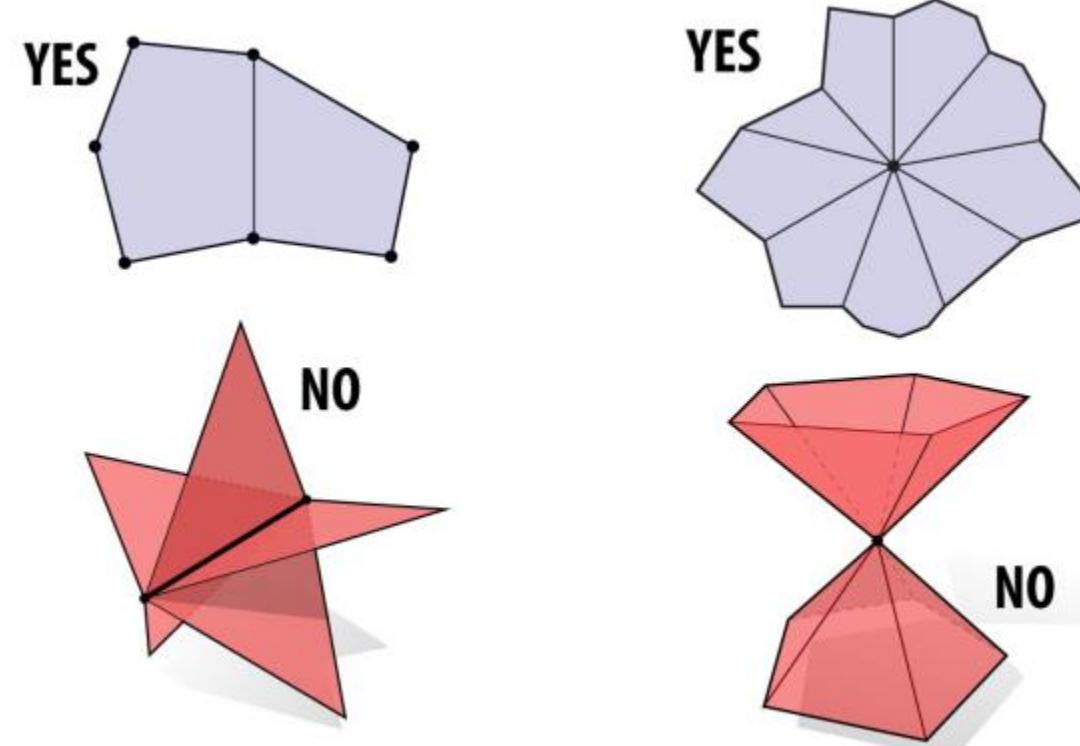


Manifold vs Non-Manifold

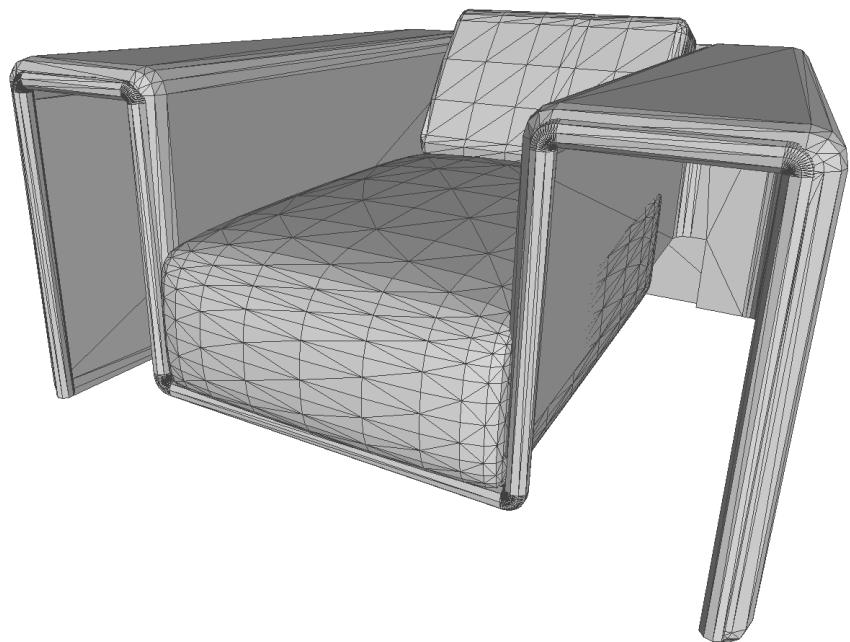
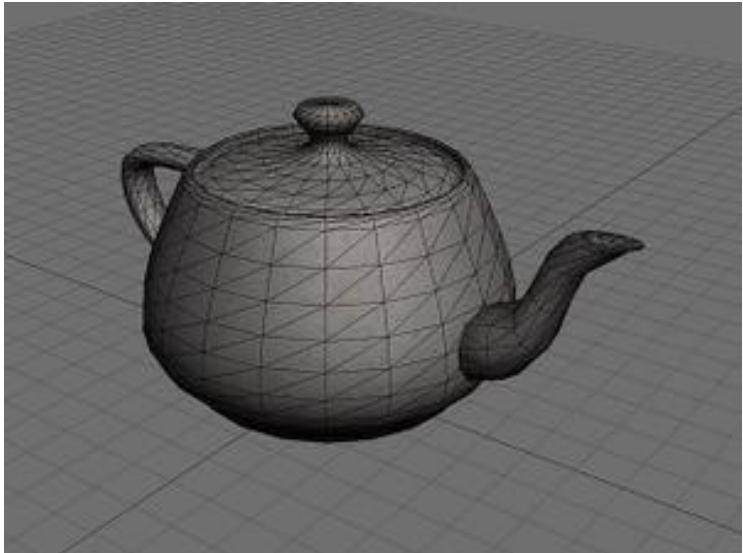


Manifold Polygon Mesh

- Each edge is contained in only two polygons
- The polygons containing each vertex make a single “fan”

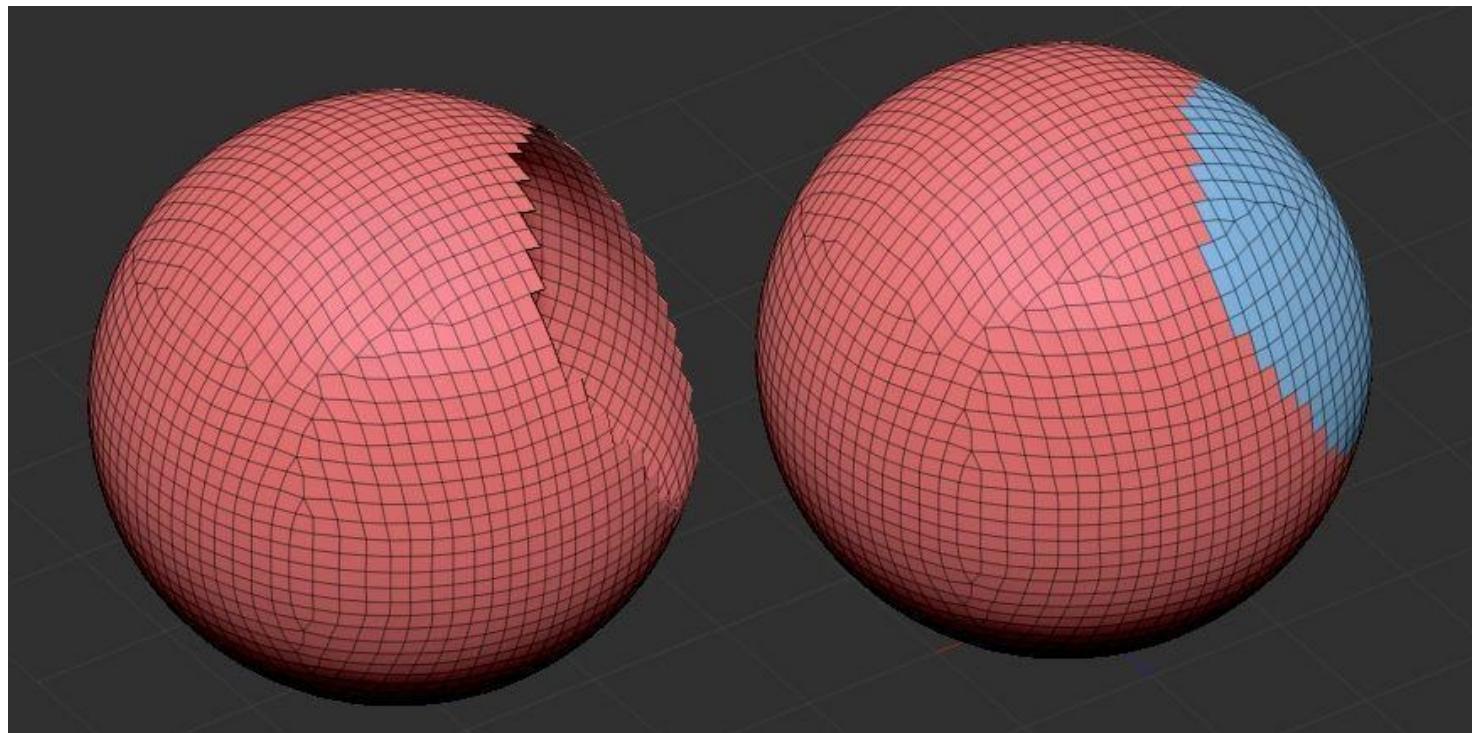


Manifold Assumption



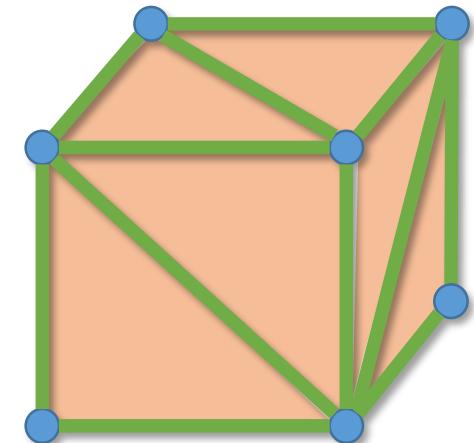
Watertight Mesh

- Does not have holes
- Has a designated inside and outside



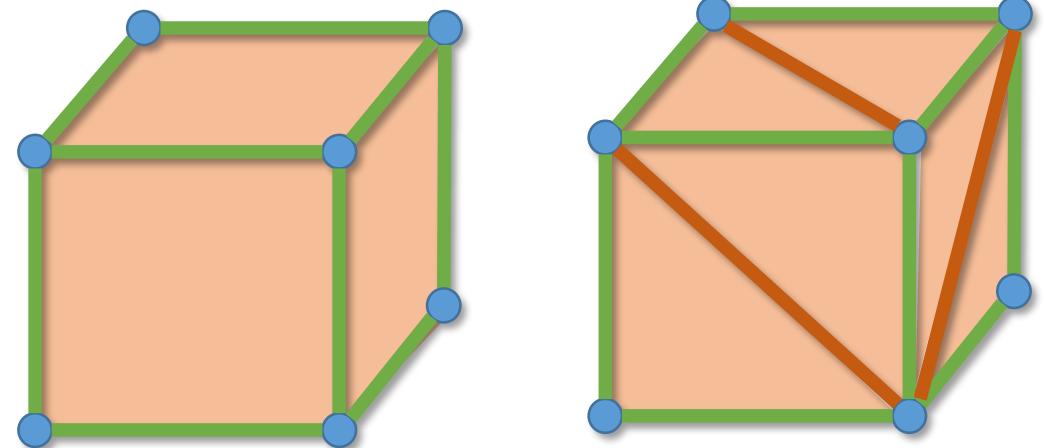
Triangle Mesh

- Polygon mesh where every face is a triangle
- Popular Representation
- Simplifies data structures
- Simplifies rendering
- Simplifies algorithms
- Each face is planar and convex
- Any polygon can be triangulated



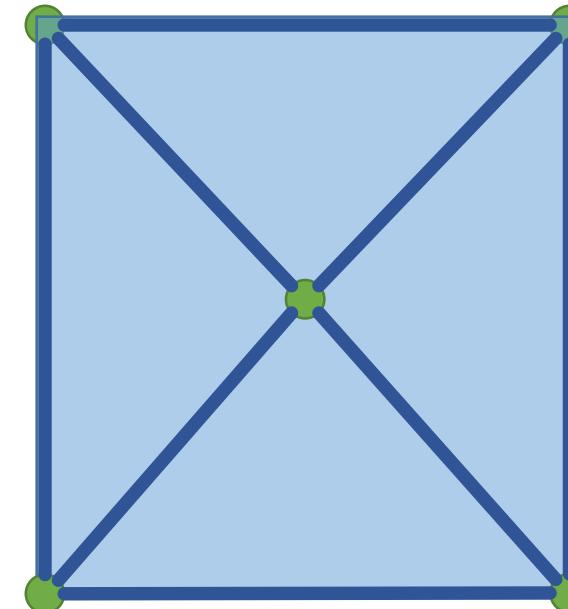
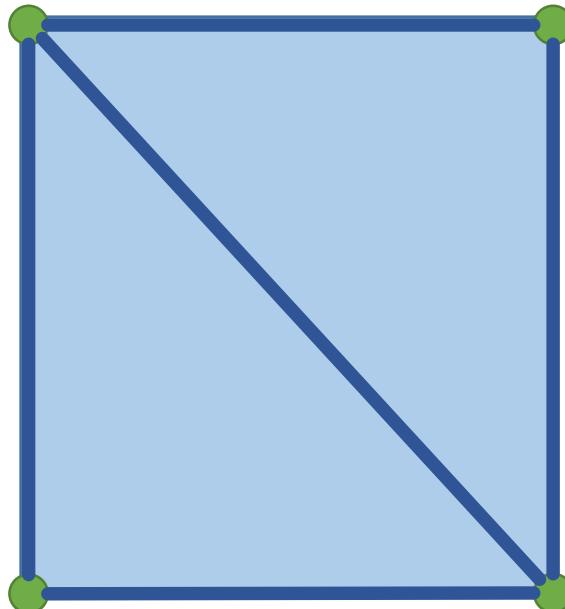
Triangle Mesh

- Polygon mesh where every face is a triangle
- Popular Representation
- Simplifies data structures
- Simplifies rendering
- Simplifies algorithms
- Each face is planar and convex
- Any polygon can be triangulated

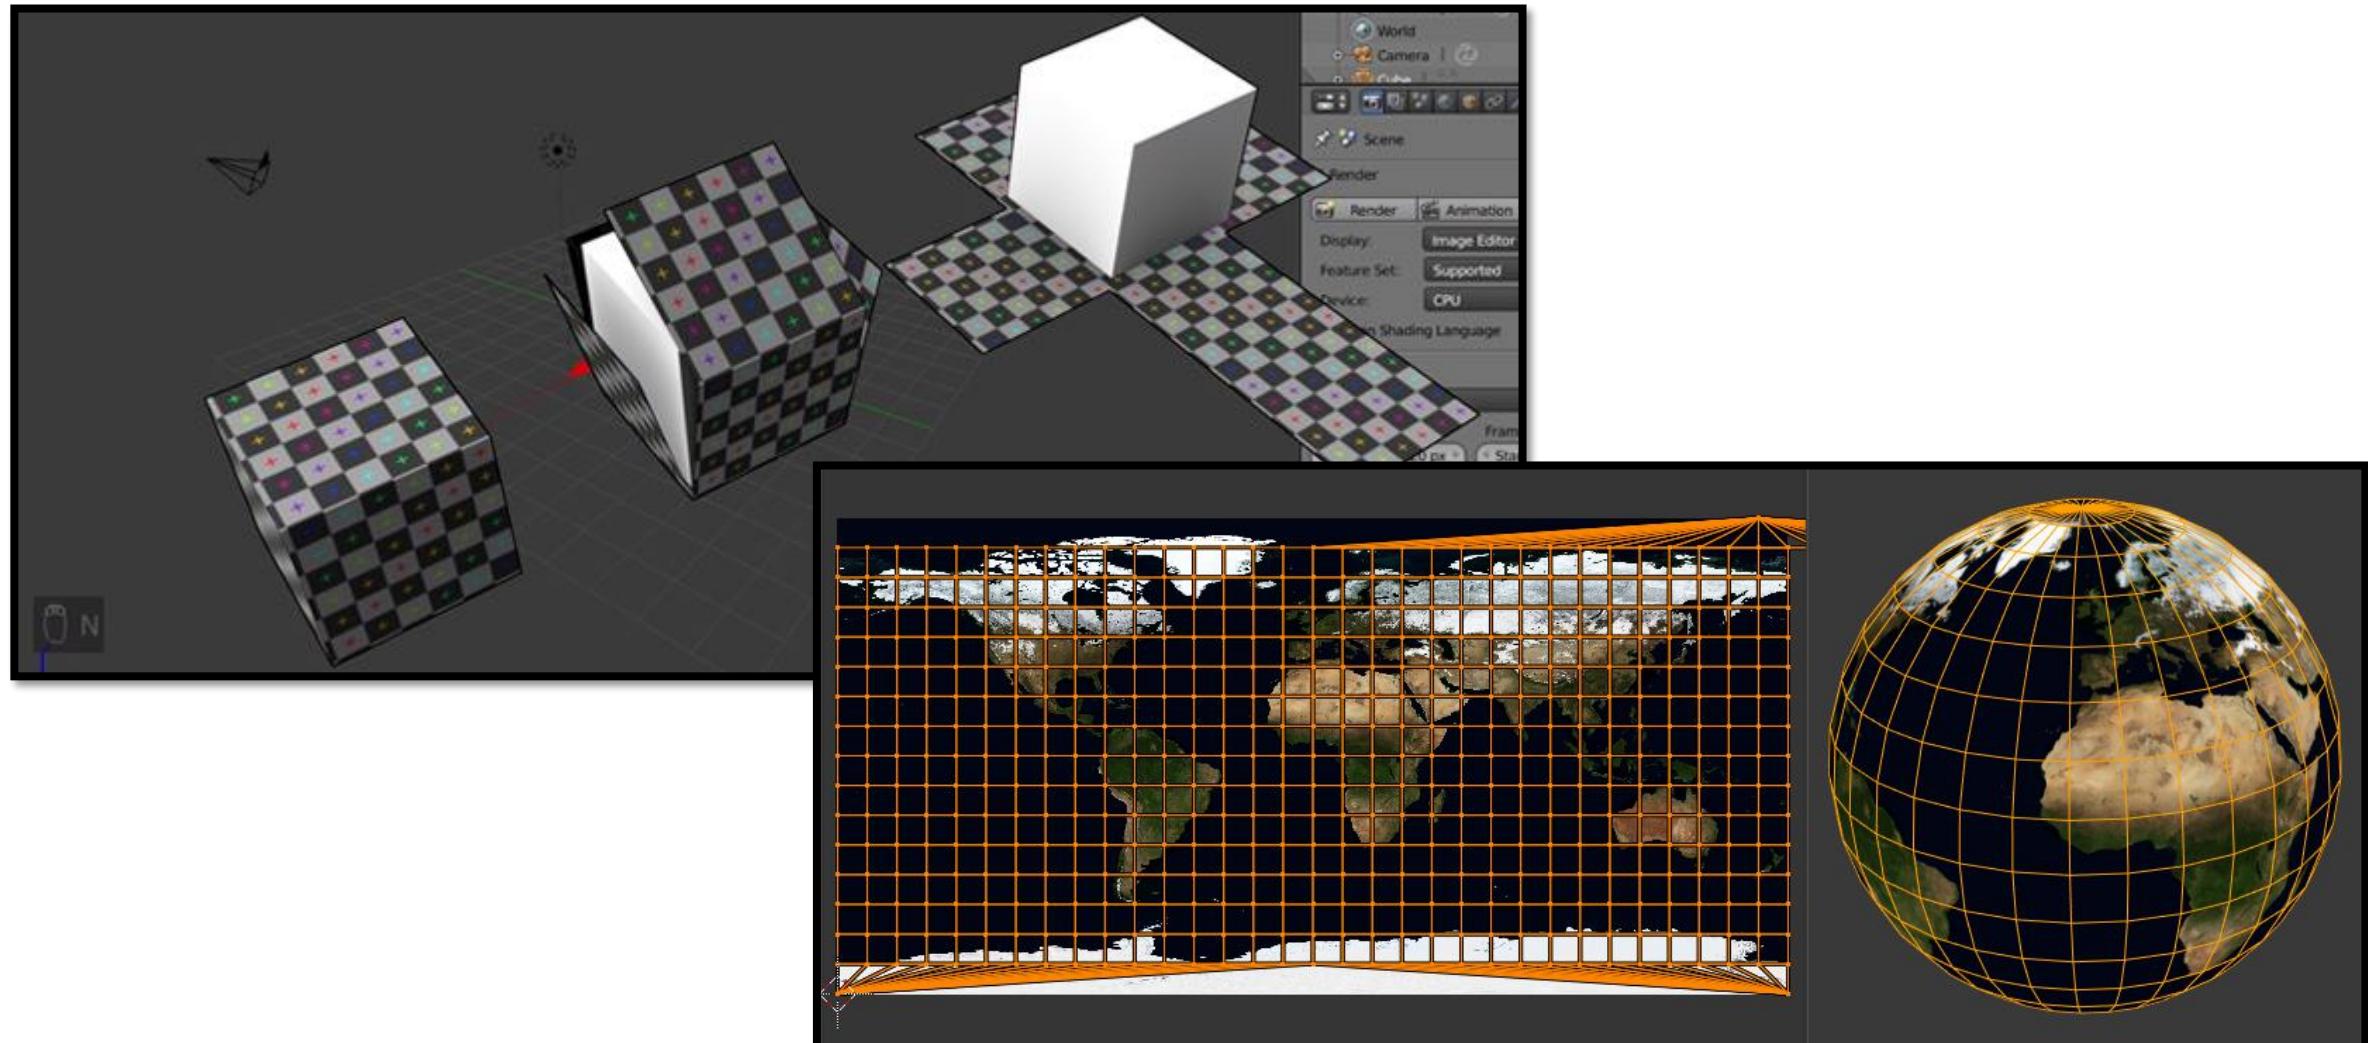


Polygon Meshes

- Many different ways to describe the same surface



Textured Polygon Meshes



Polygon Meshes – Data Structure

- What should be stored?
- Geometry: vertices, edges, faces
- Connectivity
 - adjacency relationships
- Attributes
 - normals, color, texture coordinates
 - per-vertex, per-face, per-edge



Polygon Meshes – Simple Data Structures

- Triangle List
 - Face: 3 positions
 - 4 bytes per coordinate
 - 36 bytes per face
- No connectivity information
- STL format

Triangles			
0	x0	y0	z0
1	x1	y1	z1
2	x2	y2	z2
3	x3	y3	z3
4	x4	y4	z4
5	x5	y5	z5
...

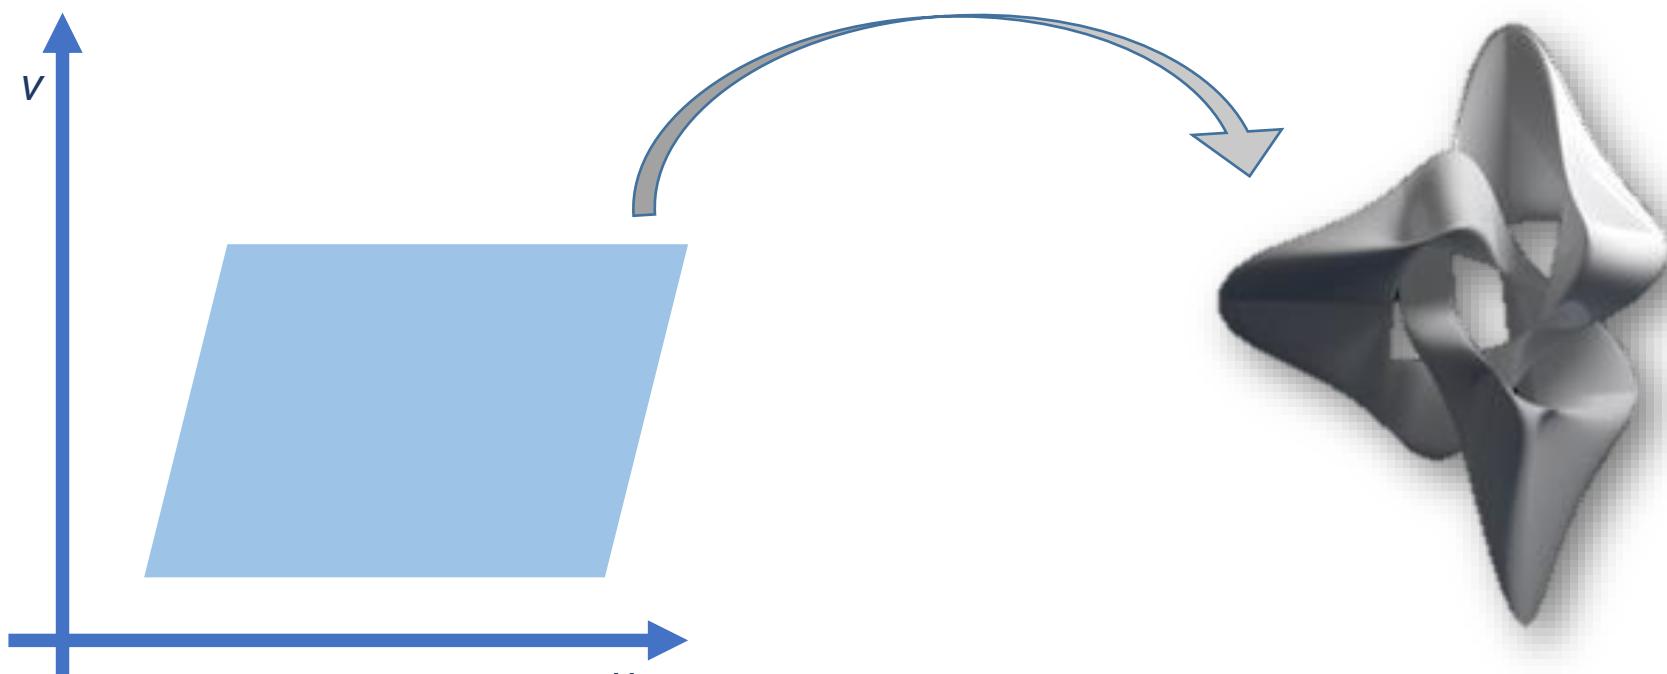
Polygon Meshes – Simple Data Structures

- Indexed Face Set
 - Vertex position
 - Face: vertex indices
 - 12 bytes per vertex
 - 12 bytes per face
- No explicit neighborhood information
- OBJ, OFF, WRL formats

Triangles			
0	v0	v1	v2
1	v0	v1	v3
2	v2	v4	v3
3	v5	v2	v6
...

Vertices			
v0	x0	y0	z0
v1	x1	y1	z1
v2	x2	y2	z2
v3	x3	y3	z3
v4	x4	y4	z4
v5	x5	y5	z5
...

Parametric Surfaces



$$\mathbf{p}: \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

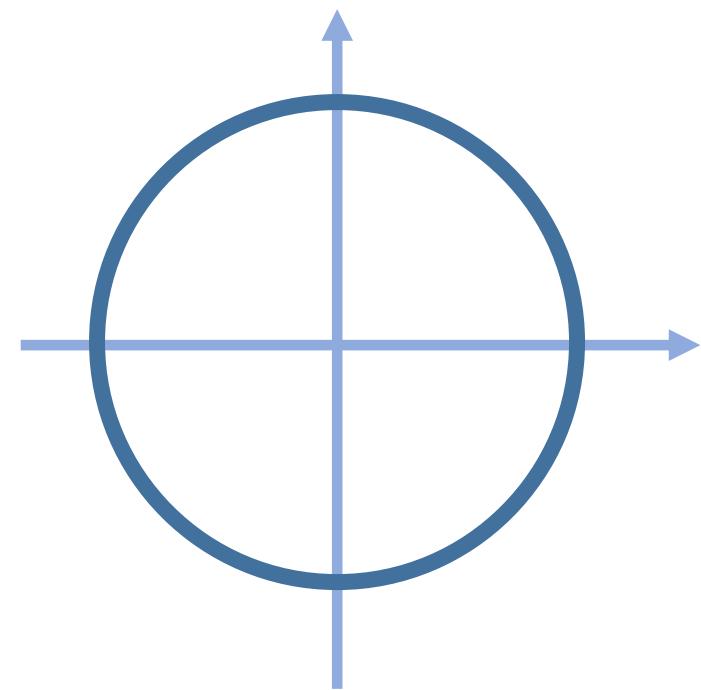
$$\mathbf{p}(u, v) = (x(u, v), y(u, v), z(u, v))$$

Parametric Curves

$$\begin{aligned} \mathbf{p}: \mathbb{R} &\rightarrow \mathbb{R}^2 \\ t \mapsto \mathbf{p}(t) &= (x(t), y(t)) \end{aligned}$$

Circle:

$$\begin{aligned} \mathbf{p}(t) &= r(\cos(t), \sin(t)) \\ t &\in [0, 2\pi) \end{aligned}$$

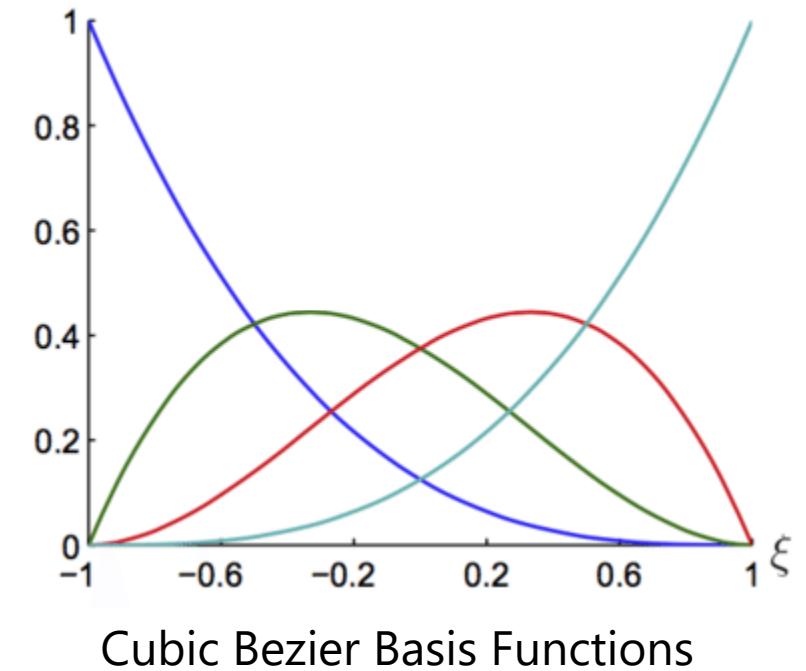
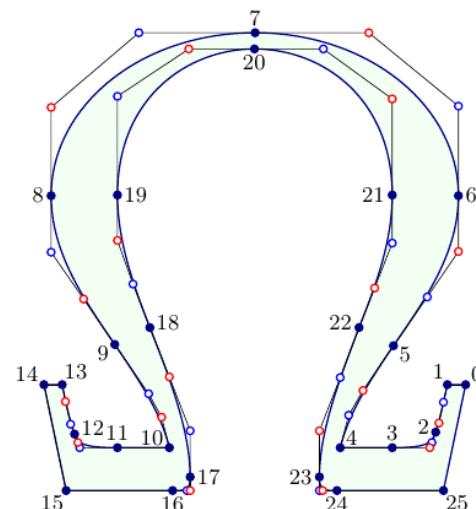
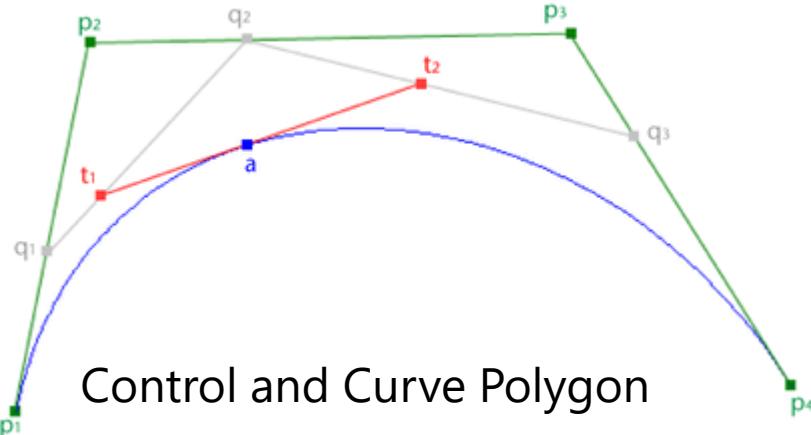


Parametric Curves

- Bezier curves, splines

$$s(t) = \sum_{i=0}^n p_i B_i^n(t)$$
$$\sum_{i=0}^n B_i^n(t) = \binom{n}{i} t^i (1-t)^{n-i}$$

Control Points



Cubic Bezier Basis Functions

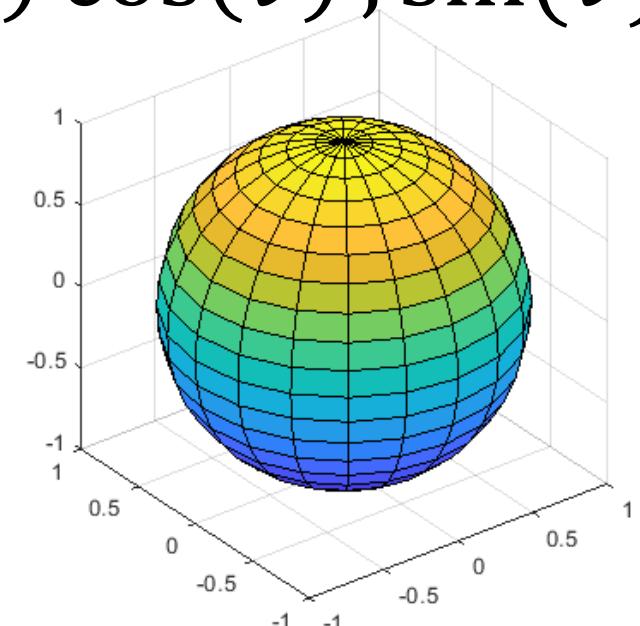
Parametric Surfaces

- Example: Sphere

$$p: \mathbb{R}^2 \rightarrow \mathbb{R}^3$$

$$p(u, v) = r(\cos(u) \cos(v), \sin(u) \cos(v), \sin(v))$$

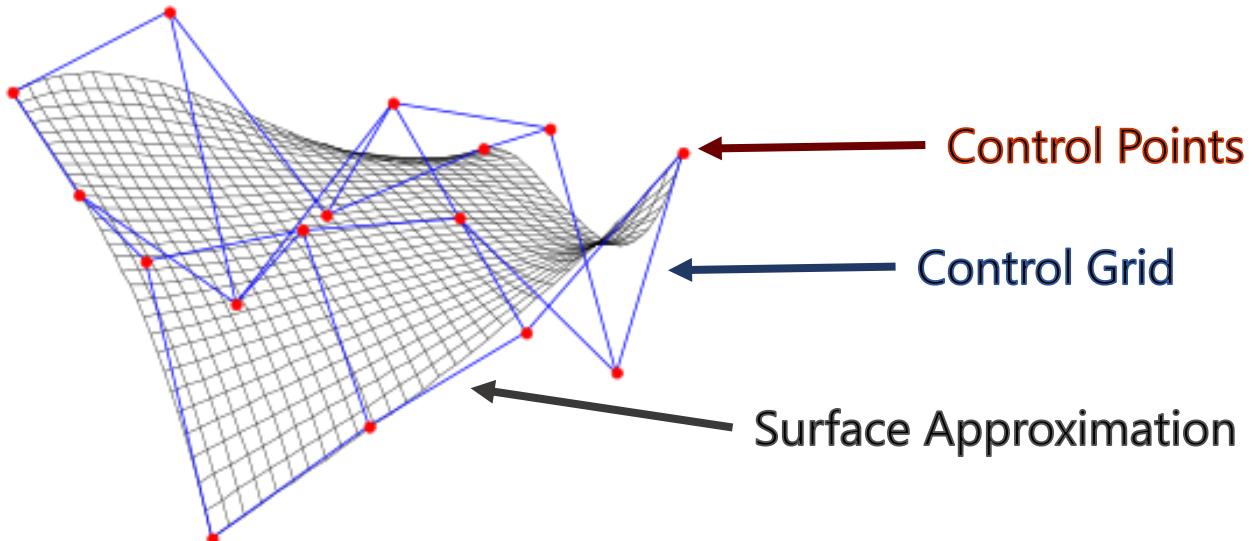
$$(u, v) \in [0, 2\pi) \times \left[-\frac{\pi}{2}, \frac{\pi}{2}\right]$$



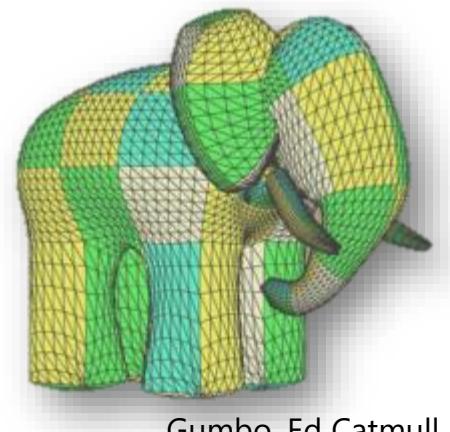
Parametric Surfaces

- Bezier Surface
 - $(n + 1)(m + 1)$ control points $p_{i,j}$

$$s(u, v) = \sum_{i=0}^n \sum_{j=0}^m p_{i,j} B_i^n(u) B_j^m(v)$$



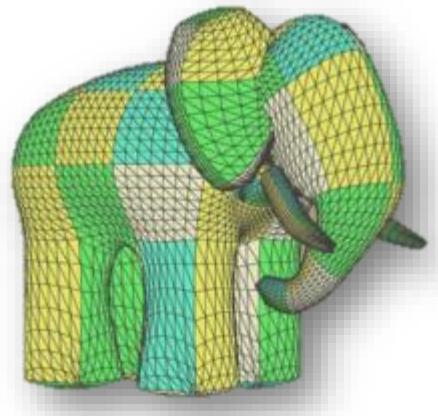
Bicubic Patches ($n=3, m=3$)



Gumbo, Ed Catmull

Parametric Surfaces

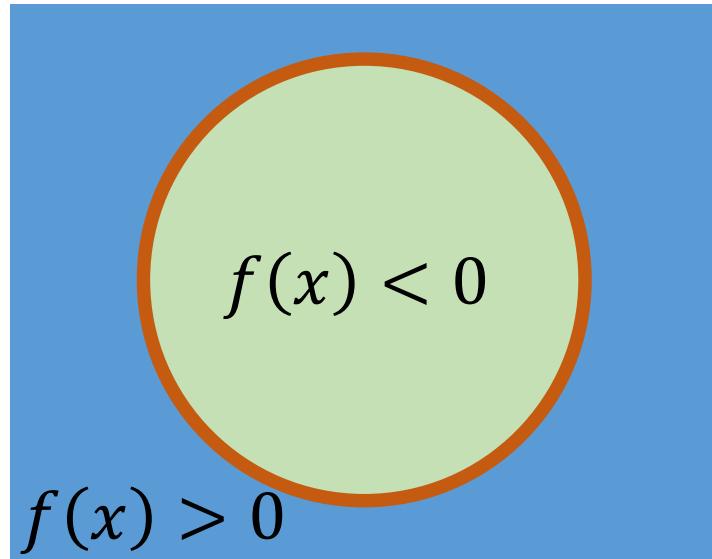
- Bezier Patch Meshes
 - Requires fewer points to represent a surface than a triangle mesh
 - Easy to manipulate
 - E.g., transforms the same way as control points under linear transformations
 - Easy to generate points on the surface
 - Easier to ensure continuity
 - Difficult to determine inside/outside
 - Difficult to determine if point is on the surface
 - Rendering more complex –
must calculate intersections with lines



Implicit Surfaces

$$f: \mathbb{R}^m \rightarrow \mathbb{R}$$

Surface in 3D:



$$S = \{x \in \mathbb{R}^3 \mid f(x) = 0\}$$

$$\{x \in \mathbb{R}^m \mid f(x) > 0\} \quad \text{Outside}$$

$$\{x \in \mathbb{R}^m \mid f(x) = 0\} \quad \text{Surface}$$

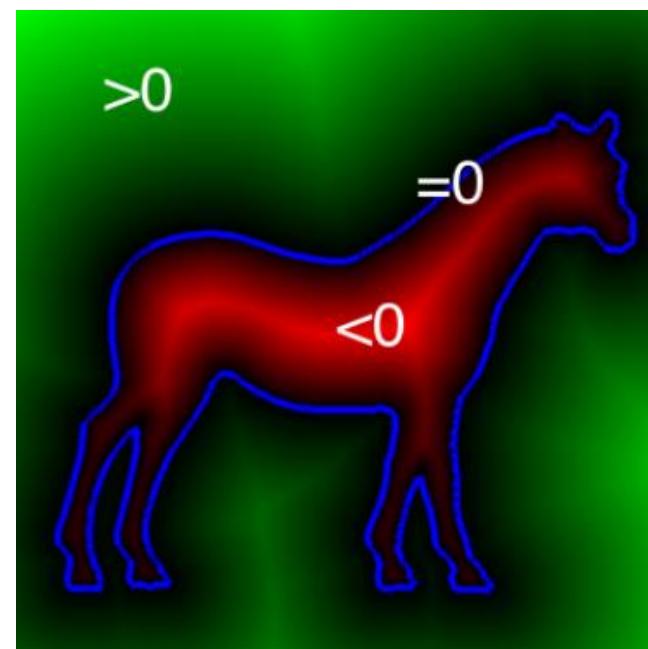
$$\{x \in \mathbb{R}^m \mid f(x) < 0\} \quad \text{Inside}$$

Implicit Surfaces

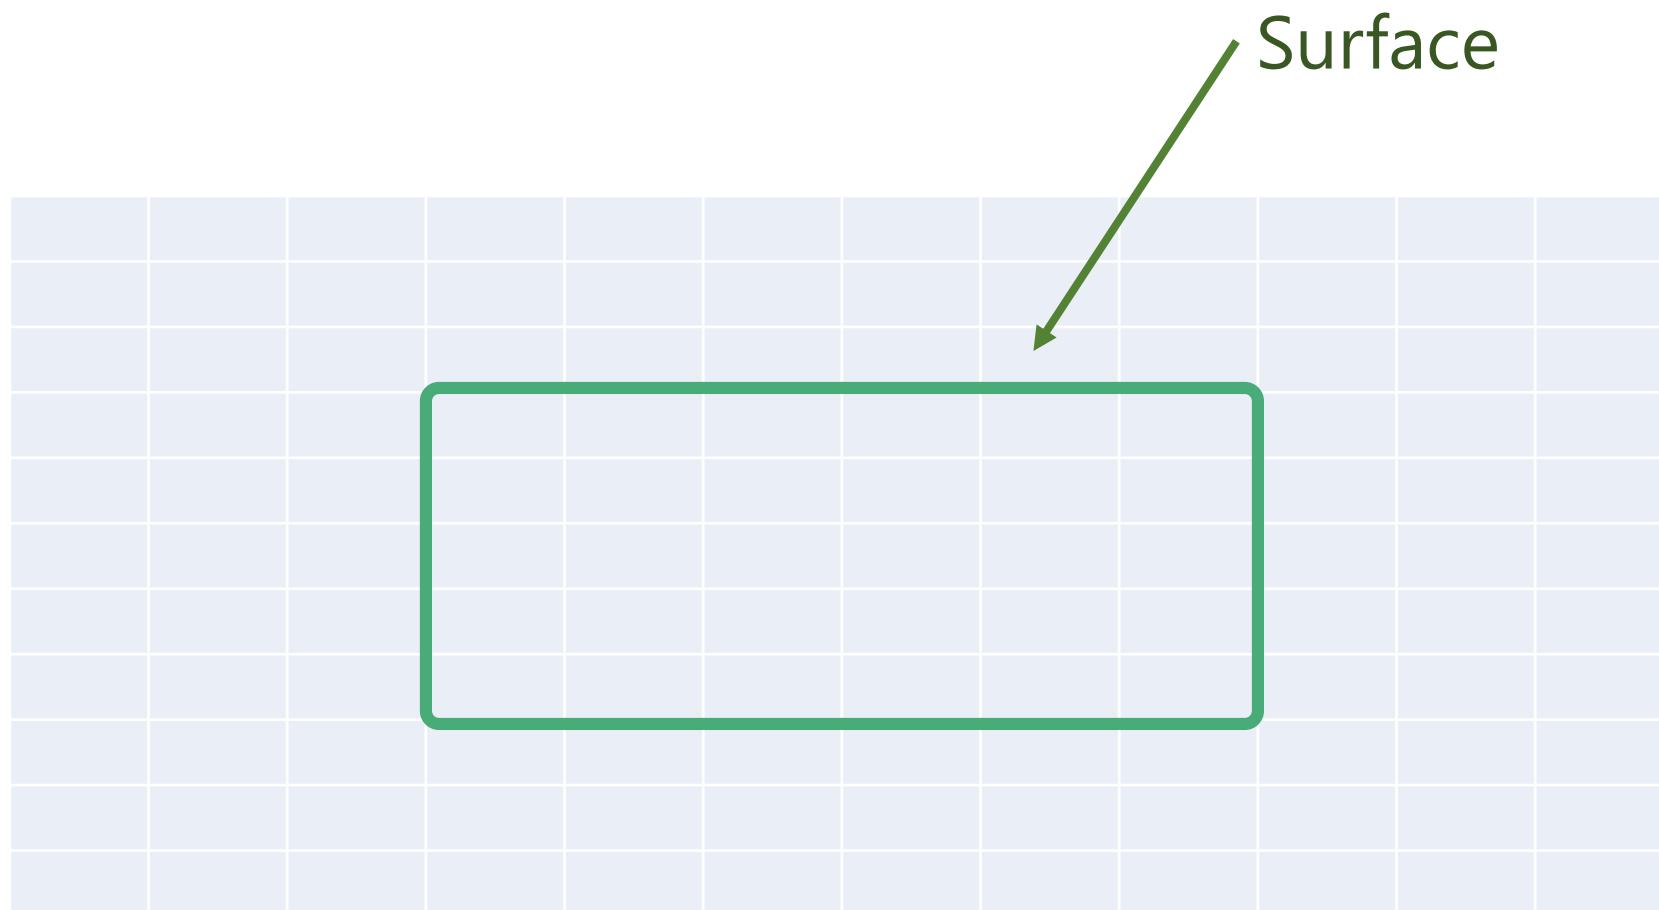
$$f: \mathbb{R}^m \rightarrow \mathbb{R}$$

Surface in 3D: $S = \{x \in \mathbb{R}^3 \mid f(x) = 0\}$

- Signed Distance Function
 - distance to the surface



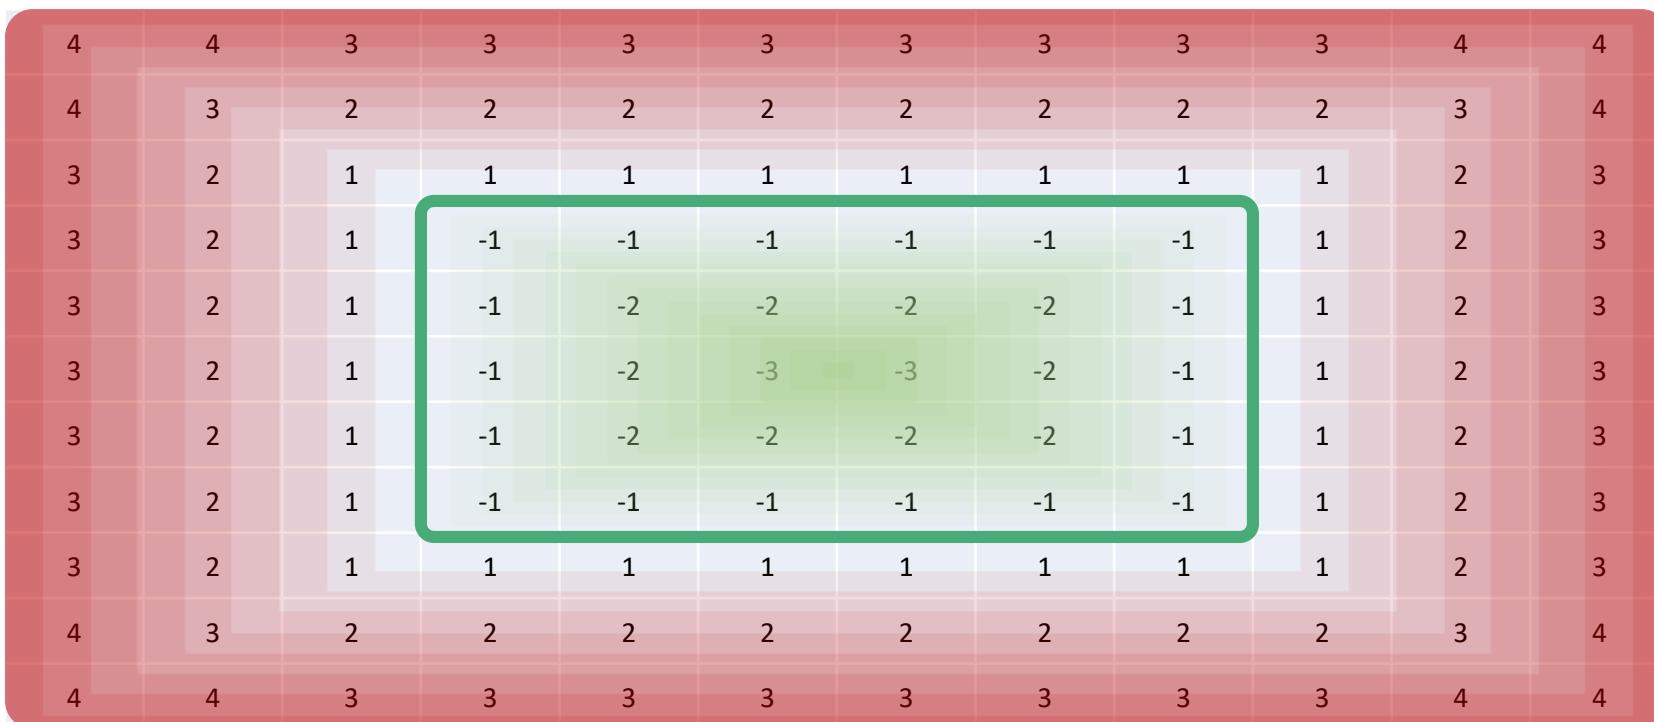
Signed Distance Field



Occupancy Grid

0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	1	1	1	1	1	1	1	1	0	0	0
0	0	1	0	0	0	0	0	0	1	0	0	0
0	0	1	0	0	0	0	0	0	1	0	0	0
0	0	1	0	0	0	0	0	0	1	0	0	0
0	0	1	0	0	0	0	0	0	1	0	0	0
0	0	1	0	0	0	0	0	0	1	0	0	0
0	0	1	1	1	1	1	1	1	1	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0

Signed Distance Field

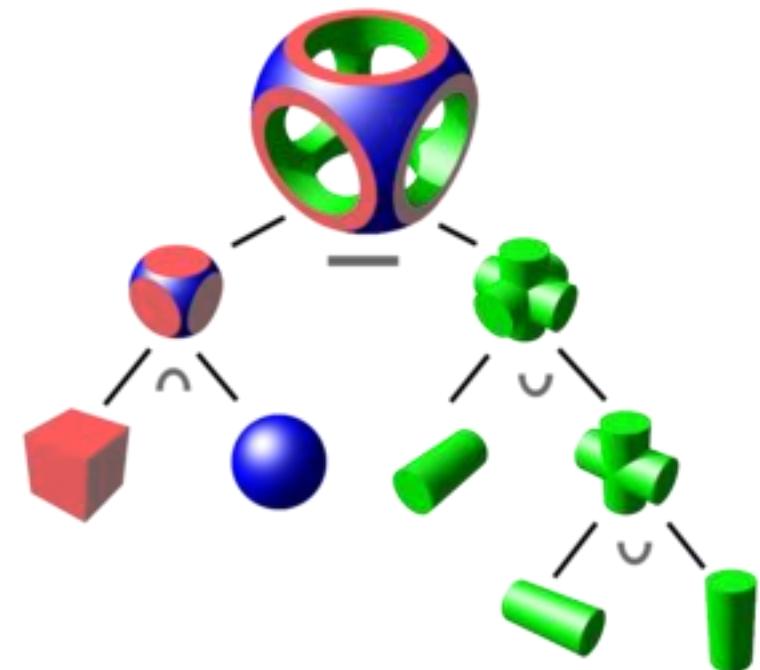
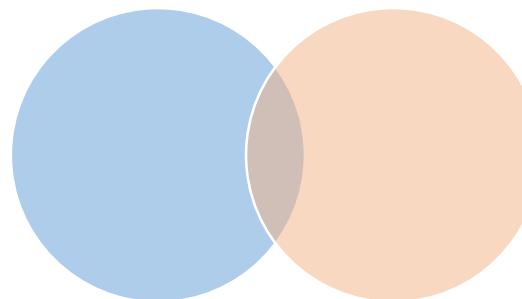


Truncated Signed Distance Field

N/A												
N/A												
N/A	N/A	1	1	1	1	1	1	1	1	N/A	N/A	N/A
N/A	N/A	1	-1	-1	-1	-1	-1	-1	1	N/A	N/A	N/A
N/A	N/A	1	-1	N/A	N/A	N/A	N/A	-1	1	N/A	N/A	N/A
N/A	N/A	1	-1	N/A	N/A	N/A	N/A	-1	1	N/A	N/A	N/A
N/A	N/A	1	-1	N/A	N/A	N/A	N/A	-1	1	N/A	N/A	N/A
N/A	N/A	1	-1	-1	-1	-1	-1	-1	1	N/A	N/A	N/A
N/A	N/A	1	1	1	1	1	1	1	1	N/A	N/A	N/A
N/A												
N/A												

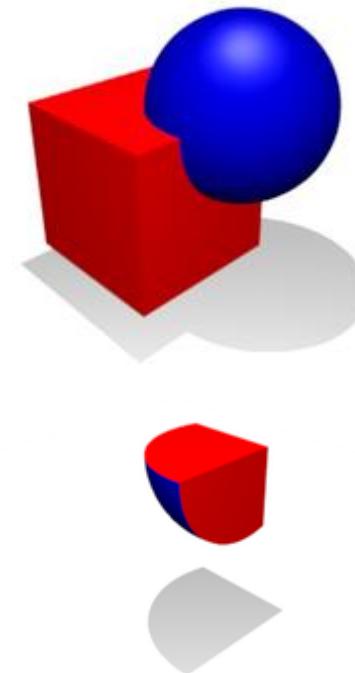
Operations on Implicit Surfaces

- Create a complex object by Boolean operators on simpler objects
- Surface is the boundary of the constructed object
- Boolean Operators (1: object, 0: not object)
 - Union
 - Intersection



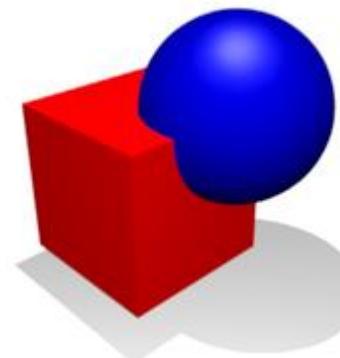
Boolean Operators on Implicit Surfaces

- Positive: outside
- Negative: inside
- Union: $\bigcup_i f_i(x) = \min f_i(x)$
- Intersection: $\bigcap_i f_i(x) = \max f_i(x)$

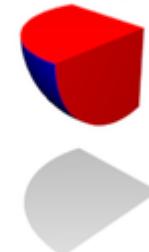


Boolean Operators on Implicit Surfaces

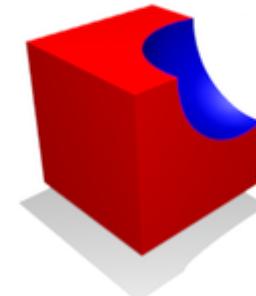
- Positive: outside
- Negative: inside
- Boolean Subtraction



Union



Intersection



Subtraction

	$f > 0$	$f < 0$
$g > 0$	$h > 0$	$h < 0$
$g < 0$	$h > 0$	$h > 0$

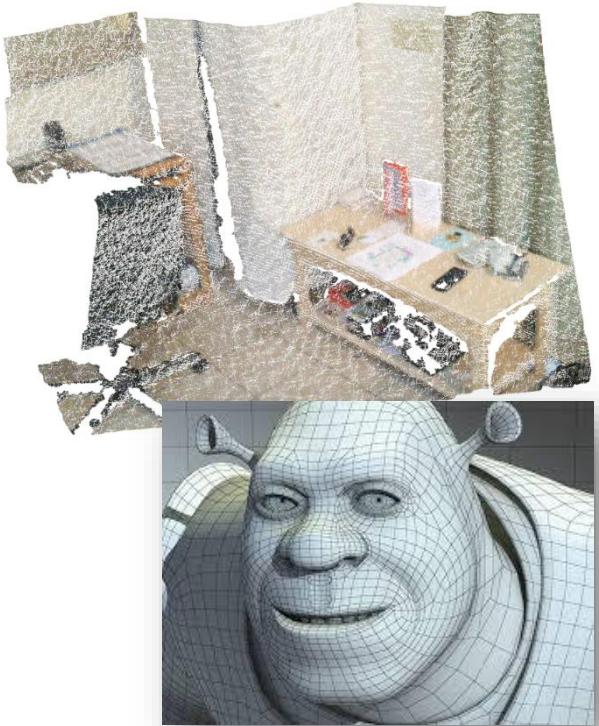
$$h = \max(f, -g)$$

- Easier than for parametric surfaces

Implicit Surfaces

- Easy to determine if a point is inside/outside
- Easy to determine if a point is on the surface
- Difficult to generate points on the surface

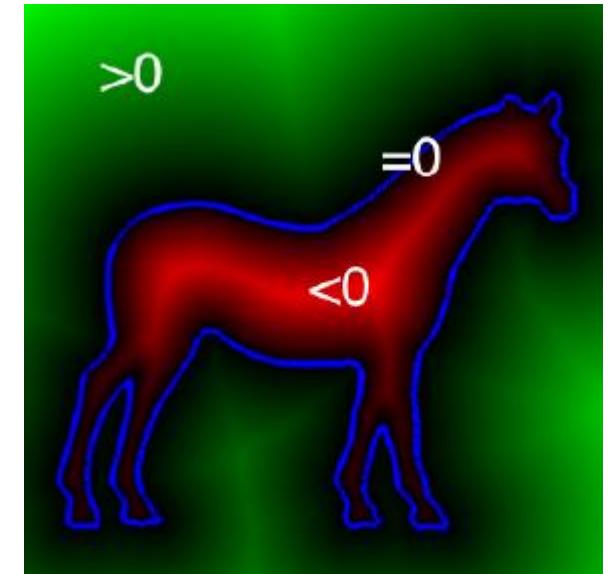
Summary



Discrete:
Meshes,
Point Samples



Parametric



Implicit:
Distance Fields

Conversion between representations

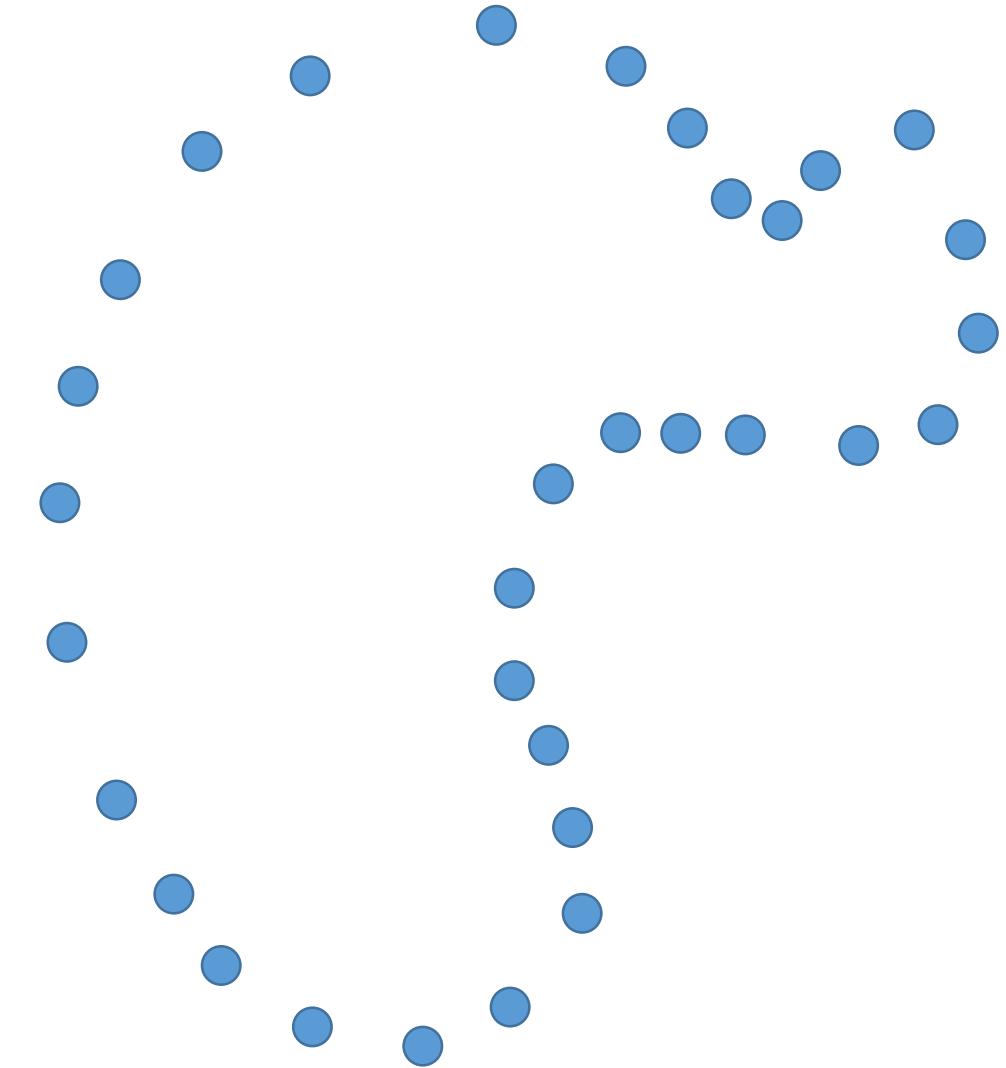


Points to Implicit

- Fit a function

$$f: \mathbb{R}^3 \rightarrow \mathbb{R}$$

with $f < 0$ outside and $f > 0$ inside

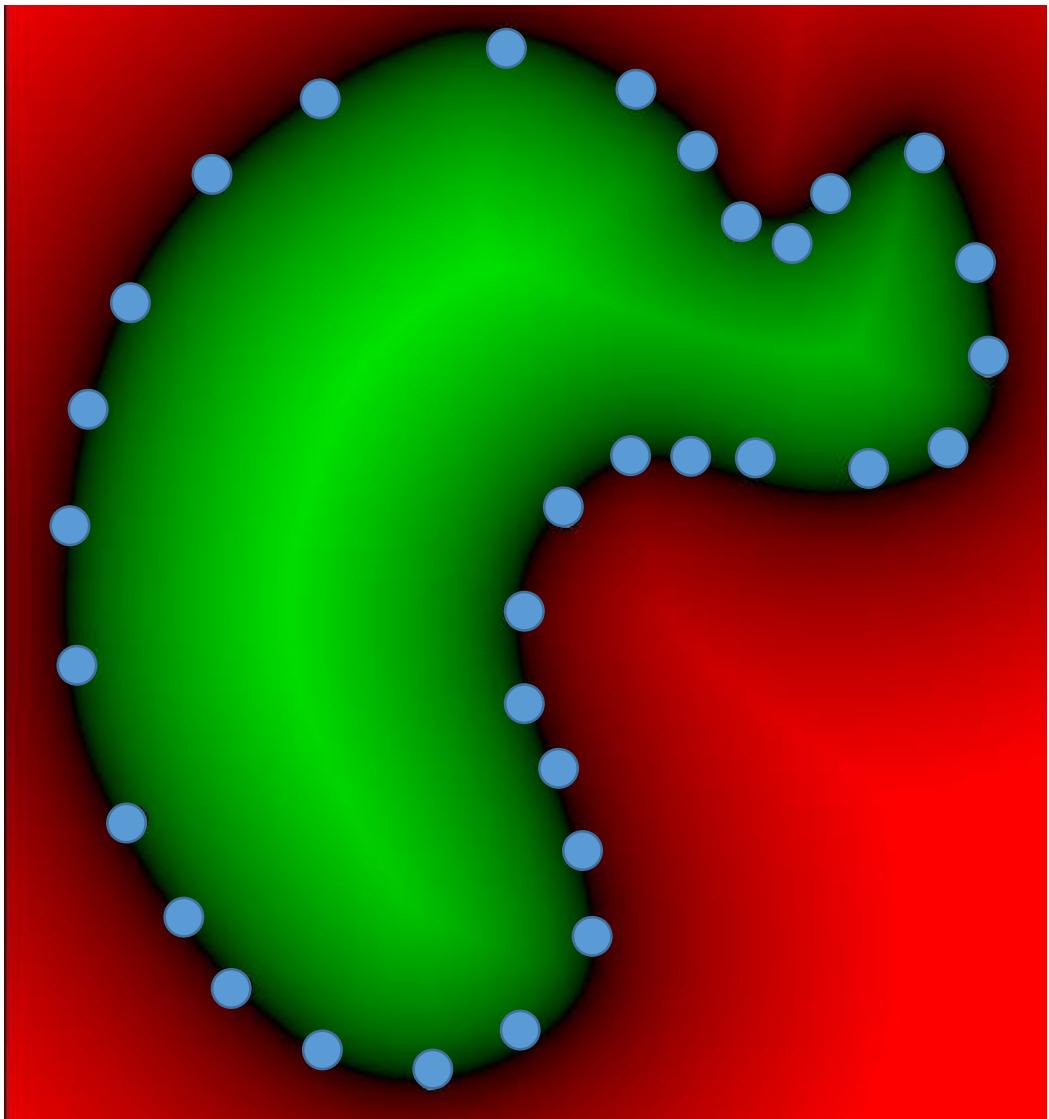
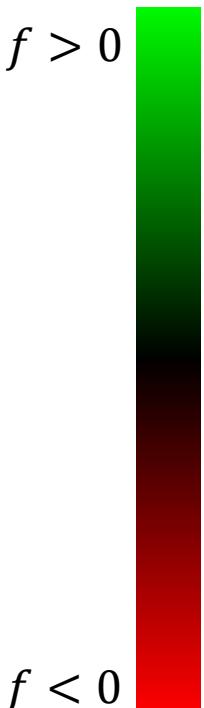


Points to Implicit

- Fit a function

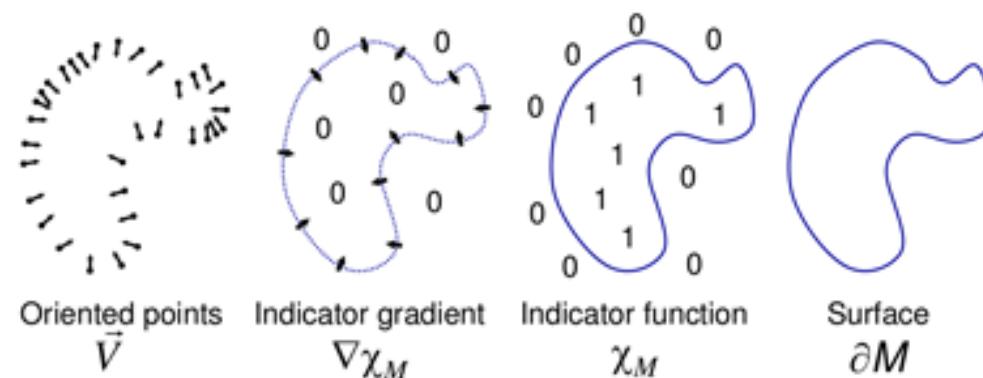
$$f: \mathbb{R}^3 \rightarrow \mathbb{R}$$

with $f < 0$ outside and $f > 0$ inside



Poisson Surface Reconstruction

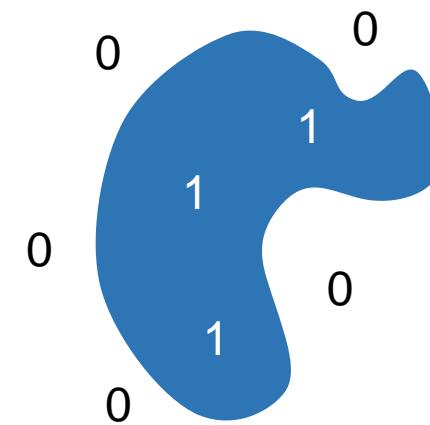
- Kazhdan et al. 2006
- Given oriented points (points with normals), compute implicit indicator function



Poisson Surface Reconstruction

- How to construct indicator function?

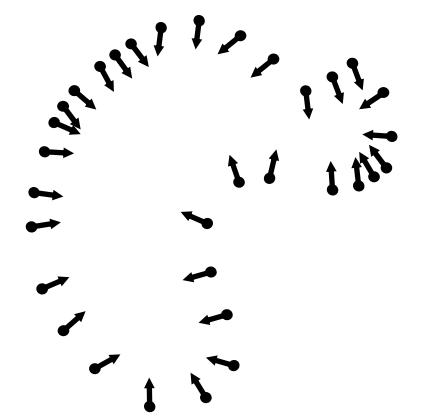
$$\chi_M(p) = \begin{cases} 1 & \text{if } p \in M \\ 0 & \text{if } p \notin M \end{cases}$$



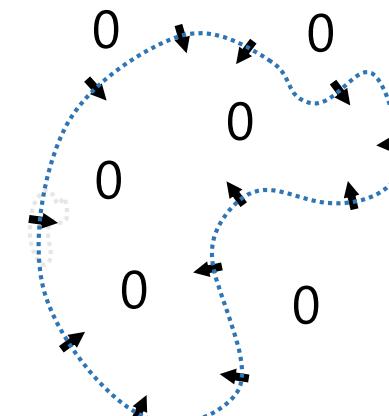
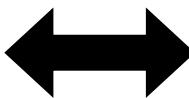
Indicator function
 χ_M

Poisson Surface Reconstruction

- Relationship between point normal and gradient of indicator function



Oriented points



Indicator gradient
 $\nabla \chi_M$

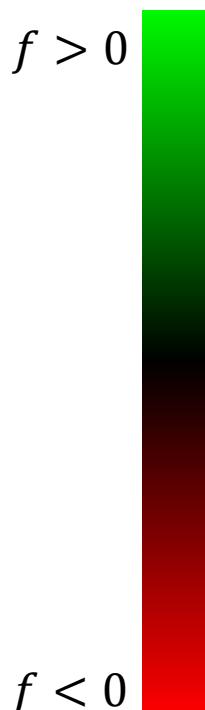
Poisson Surface Reconstruction

- Represent points by a vector field V
- Find function f whose gradient approximates V
$$\min_f \|\nabla f - V\|$$
- Transform to Poisson problem:
$$\nabla \cdot (\nabla f) = \nabla \cdot V \quad \leftrightarrow \quad \Delta f = \nabla \cdot V$$
- Solve as least squares fitting problem

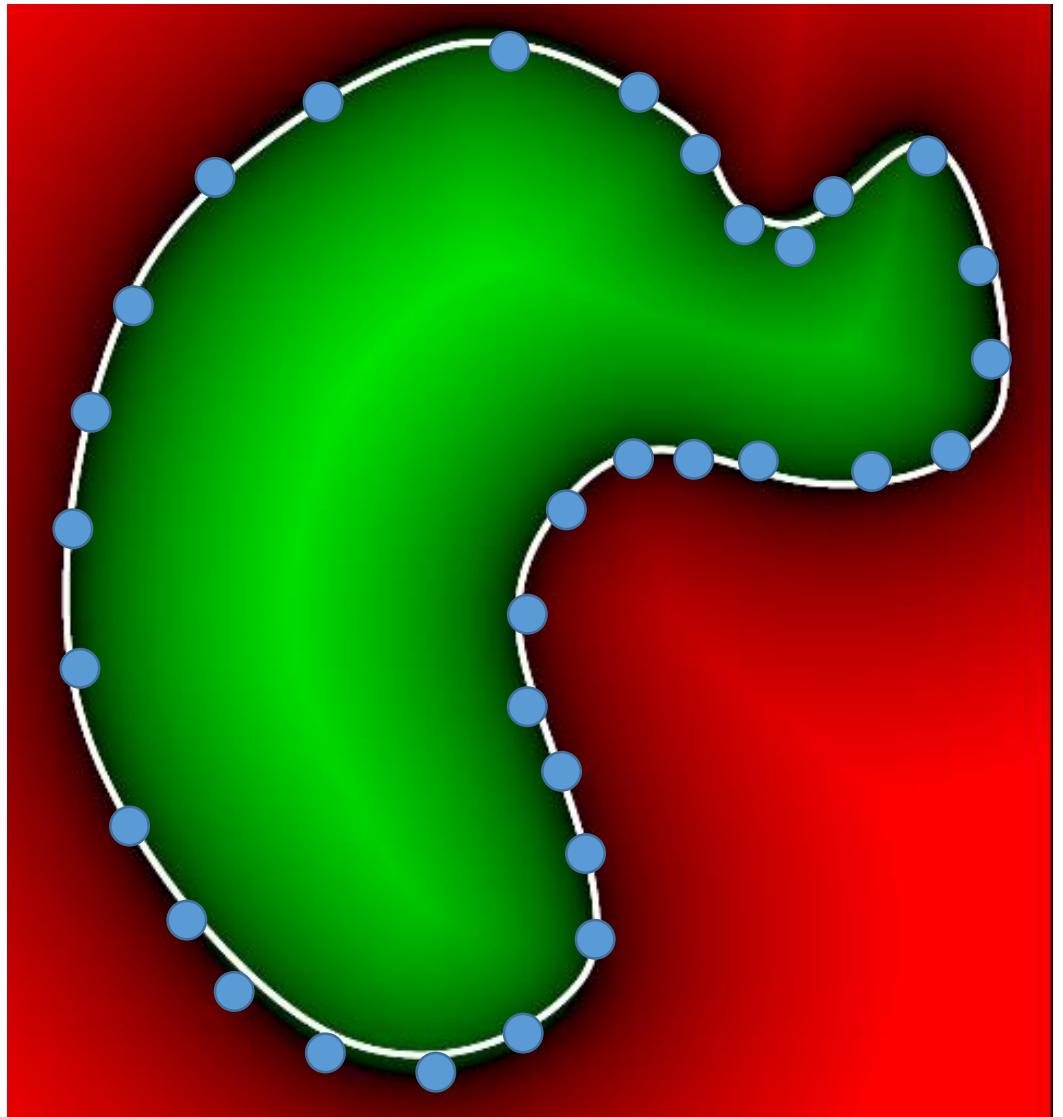
Points to Implicit to Mesh

- Fit a function
 $f: \mathbb{R}^3 \rightarrow \mathbb{R}$

with $f < 0$ outside and $f > 0$ inside

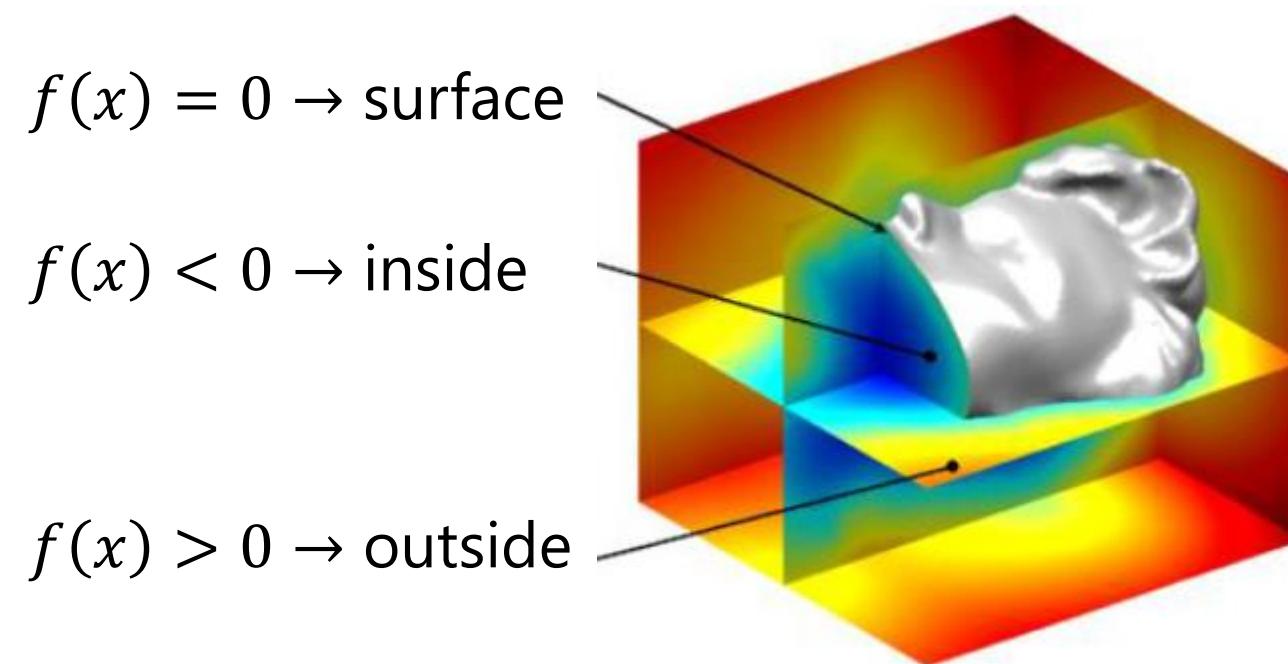


- To Mesh: Extract surface at zero-set $\{x: f(x) = 0\}$



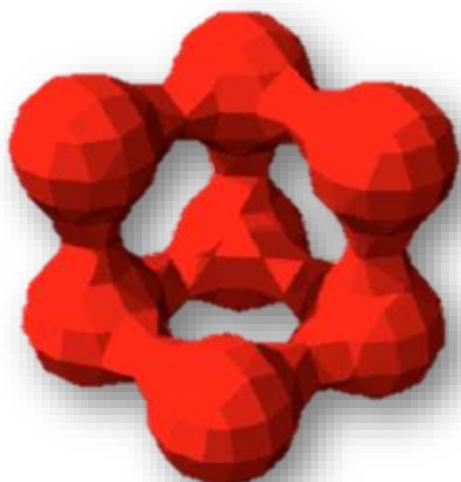
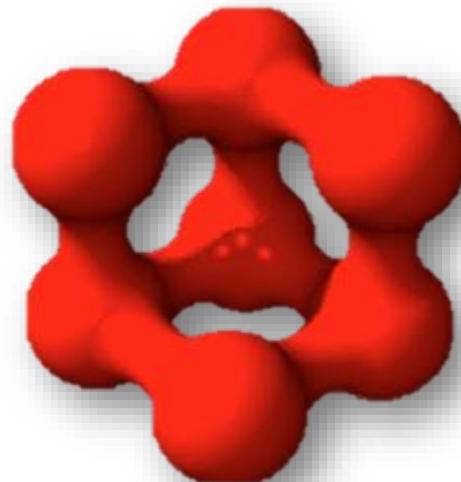
Implicit to Mesh

- Marching Cubes, Lorensen and Cline 1987
- Compute manifold mesh of level set



Marching Cubes

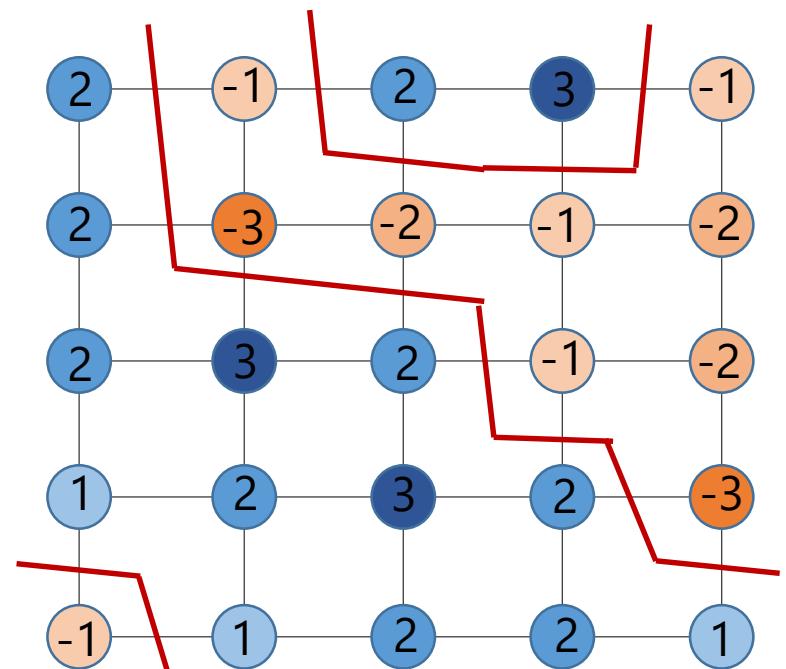
- Convert from implicit to explicit
- Given an implicit representation $\{x, \text{s.t. } f(x) = 0\}$, create a triangle mesh approximating the surface



Marching Squares (2D)

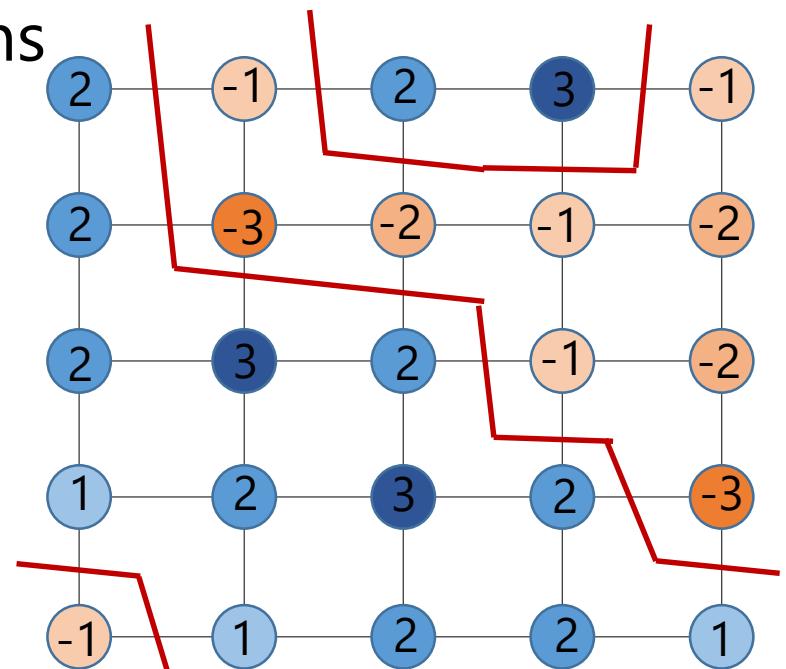
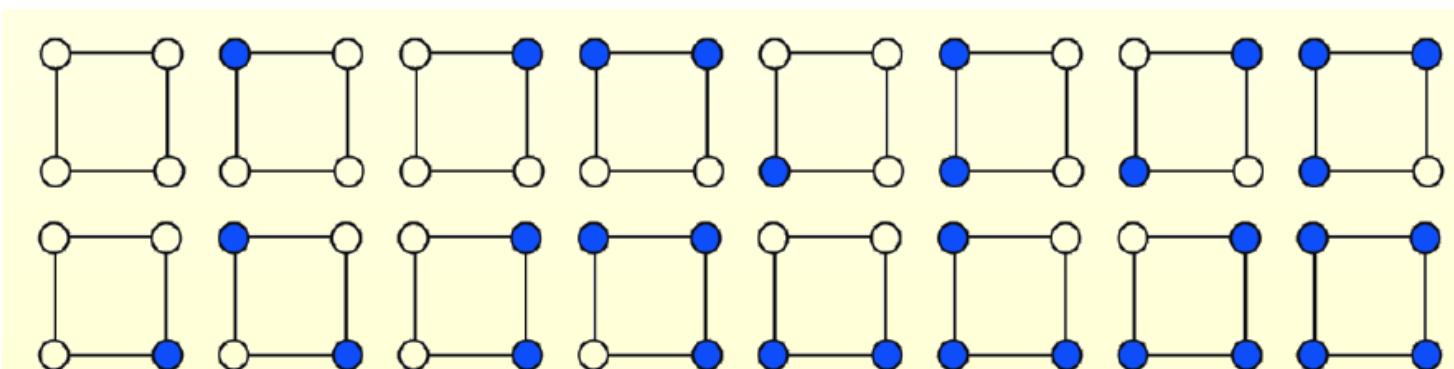
- Given: uniform sampling on a 2D grid of an implicit function f
- Determine zero-crossings: edges with a sign switch
$$f(x_1) < 0, f(x_2) > 0 \Rightarrow f(x_1 + t(x_2 - x_1)) = 0$$
for some $0 \leq t \leq 1$
- Simplest way to compute t : assume f is linear between x_1, x_2

$$t = \frac{f(x_1)}{f(x_2) - f(x_1)}$$



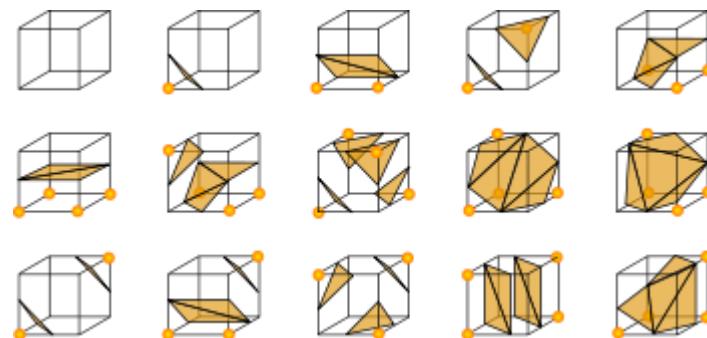
Marching Squares (2D)

- Given: uniform sampling on a 2D grid of an implicit function f
- Determine zero-crossings
- Consider all possible inside/outside combinations
- Group to lookup table

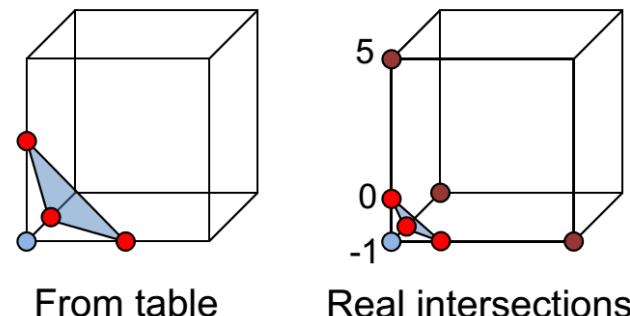


Marching Cubes (3D)

- Marching cubes lookup table

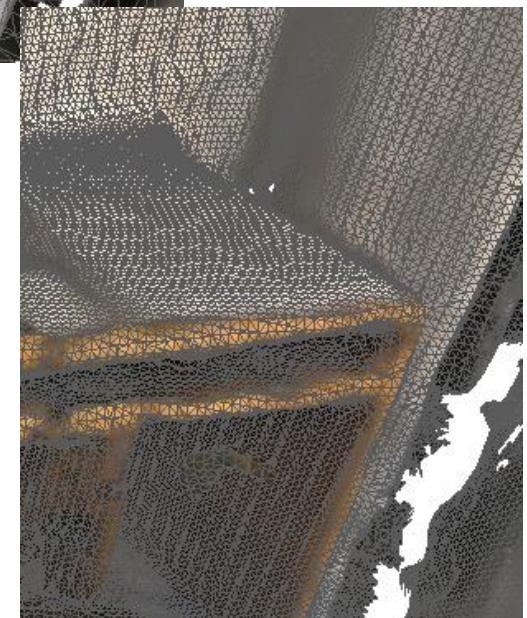
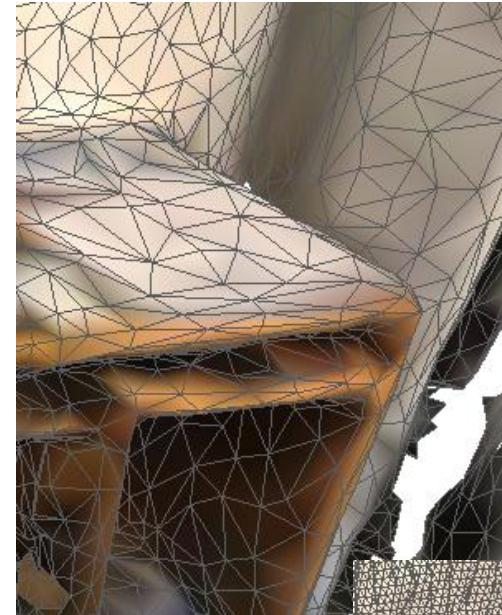


- Linear interpolation based on function values



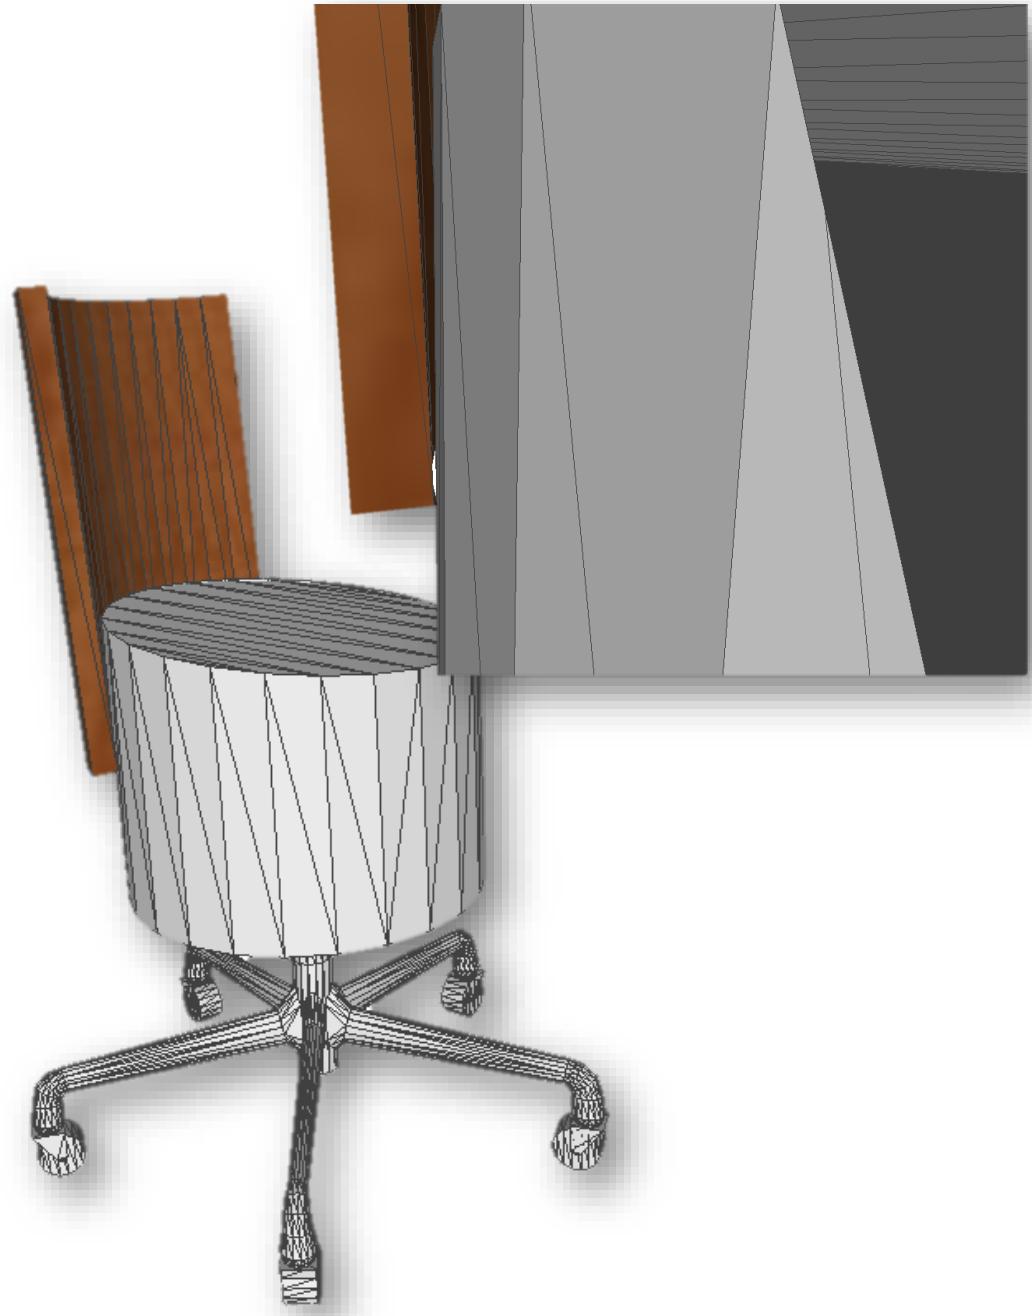
Marching Cubes

- Multi-purpose, widely applicable
 - Easy to compute and parallelizable
 - Relatively simple to implement
-
- Can create badly shaped triangles (skinny)
 - Many special cases (lookup table)
 - No sharp features



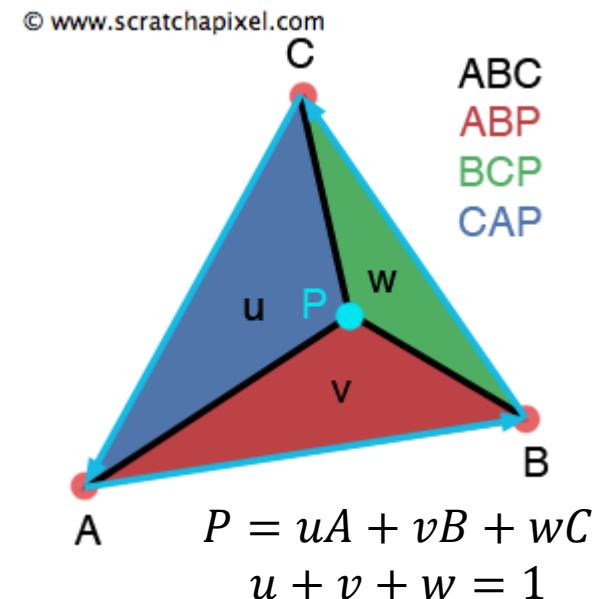
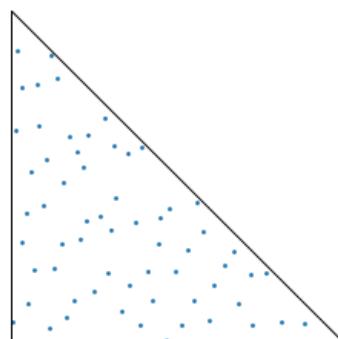
Mesh to Points

- Why?
- Some meshes can have bad triangles, bad connectivity
 - e.g., often CAD models, amateur models
- Machine learning on points
- How to sample the mesh?



Mesh to Points

- Simple approximation for uniform sampling
- Sample each triangle based on its surface area
- Sample each triangle uniformly: barycentric coordinates
- Sample random $r_1, r_2 \in [0,1]$:
 - $p = (1 - \sqrt{r_1})A + \sqrt{r_1}(1 - r_2)B + \sqrt{r_1}r_2C$

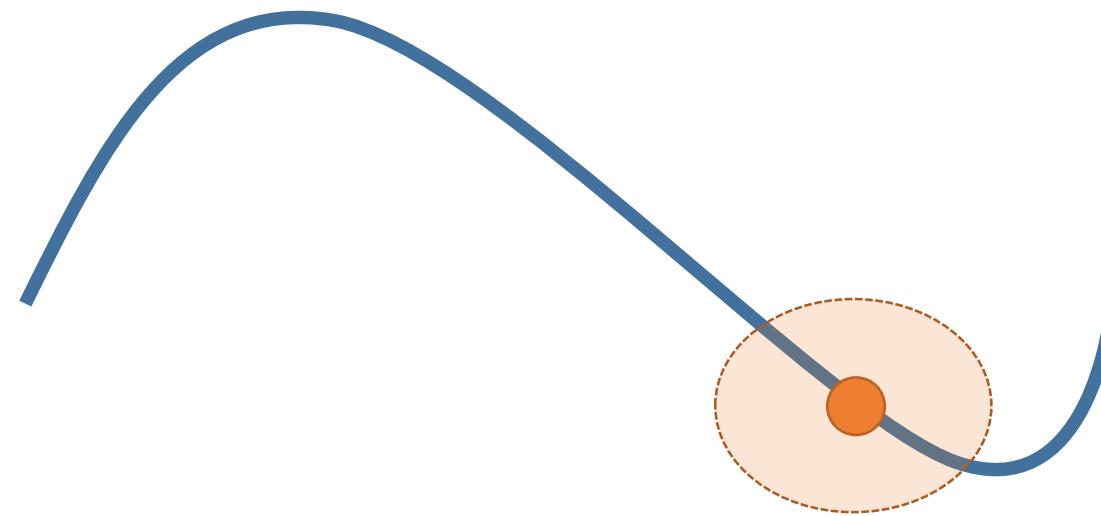


Mesh to Points

- Alternative: Farthest Point Sampling
- From an initial sample point set S , sample new point by finding the farthest point from all points in S . Iterate.
- Depends on notion of distance in domain
- Can weight adaptively (weighted distance)
- On mesh: discrete geodesic distance (path along edges)

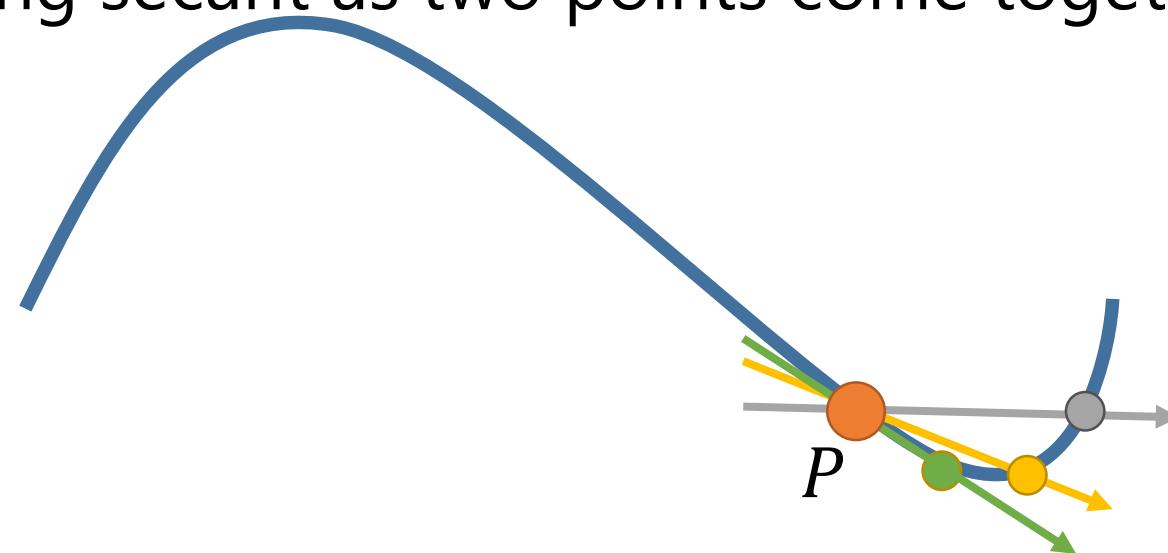
Geometric Foundations

- How to describe the geometry of a local observation: point and its neighborhood



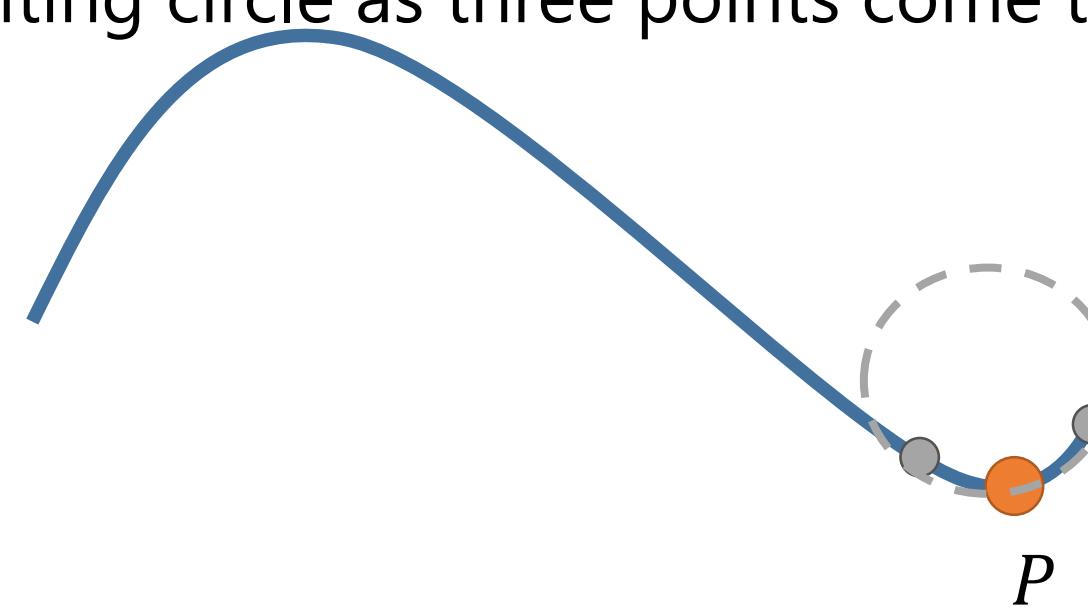
Geometric Foundations

- How to describe the geometry of a local observation: point and its neighborhood
- Tangent: limiting secant as two points come together

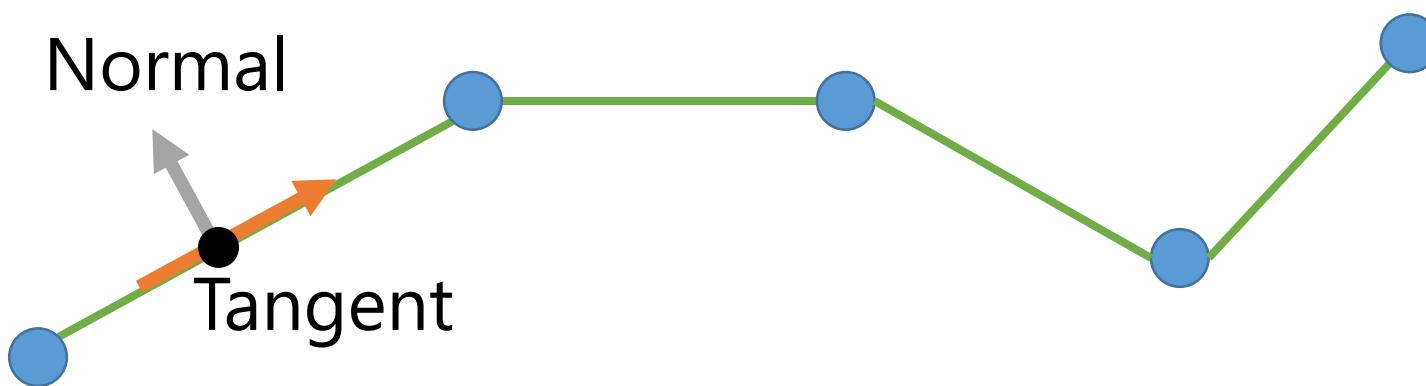


Geometric Foundations

- How to describe the geometry of a local observation: point and its neighborhood
- Curvature: limiting circle as three points come together

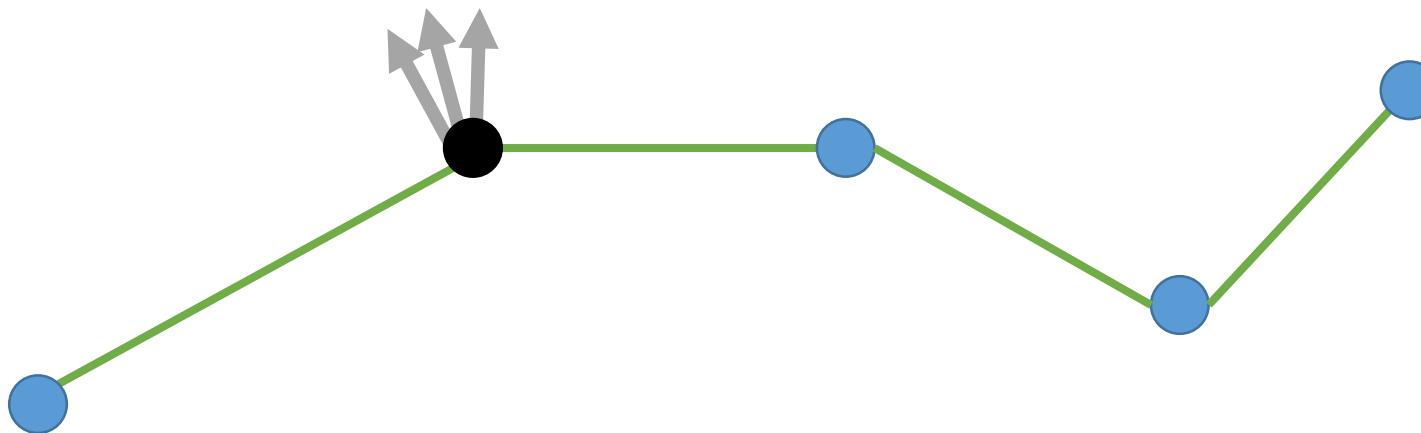


On Discrete Curves



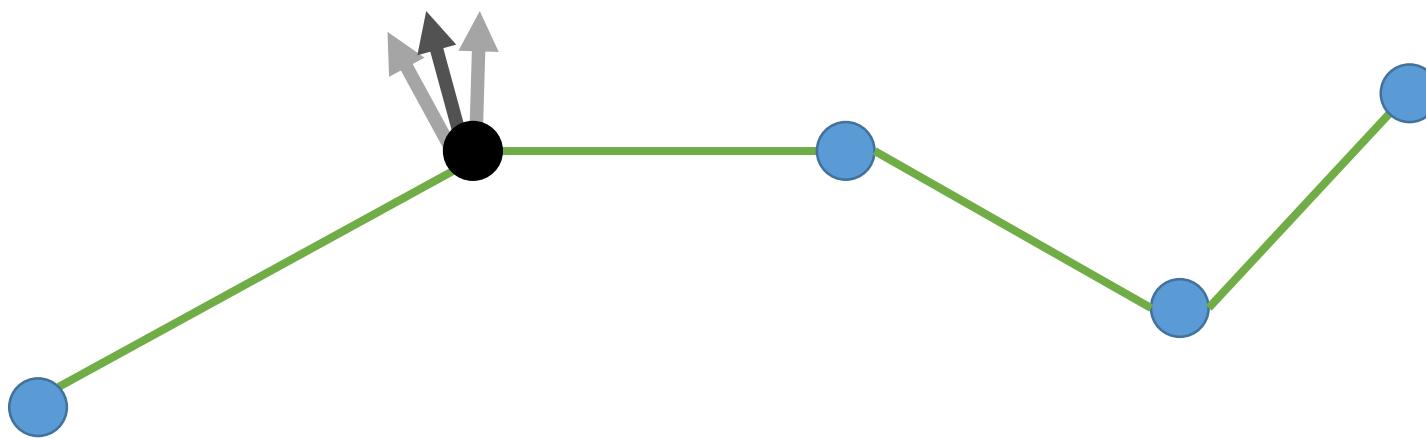
On Discrete Curves

- What about at vertices?



On Discrete Curves

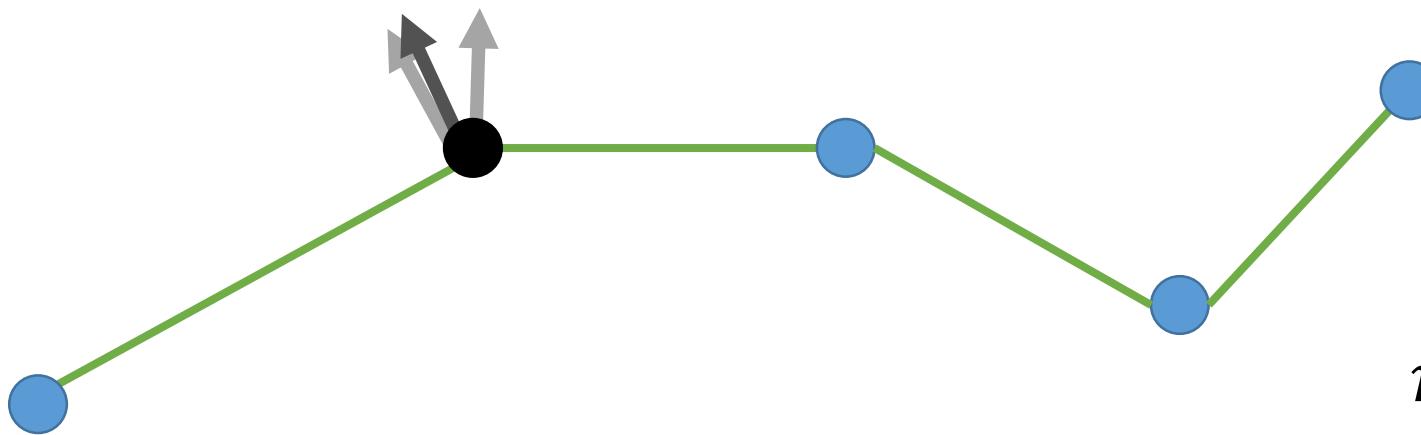
- What about at vertices?
- Can average adjacent normals



$$n_v = \frac{n_{e_1} + n_{e_2}}{\|n_{e_1} + n_{e_2}\|}$$

On Discrete Curves

- What about at vertices?
- Can average adjacent normal

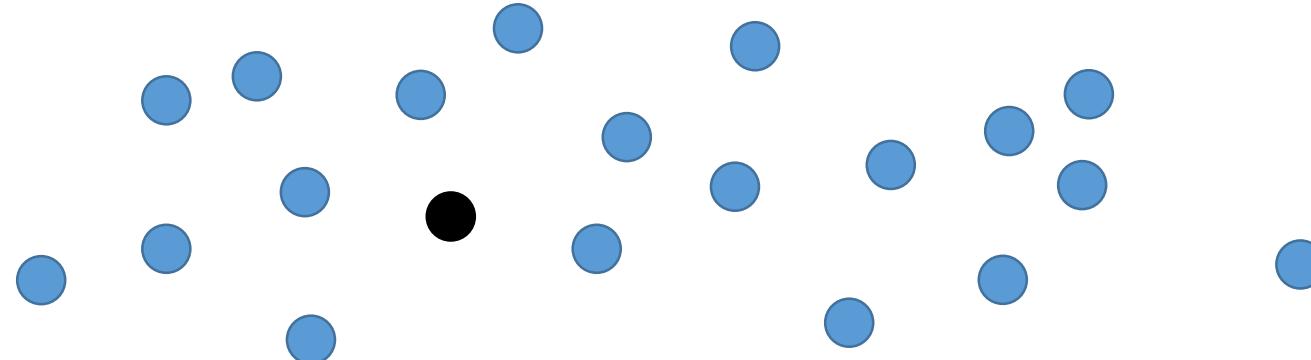


$$n_v = \frac{|e_1|n_{e_1} + |e_2|n_{e_2}}{\| |e_1|n_{e_1} + |e_2|n_{e_2} \|}$$

- Weight by edge lengths

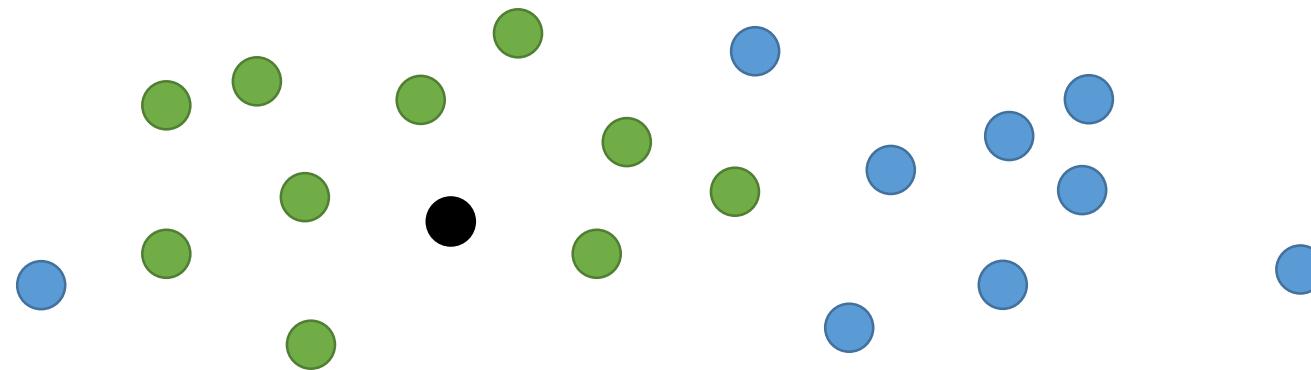
On Point Clouds

- Estimate normal by approximating the plane tangent to the surface
- Least-squares fitting problem



On Point Clouds

- Estimate normal by approximating the plane tangent to the surface
- Least-squares fitting problem
- Find neighborhood around point



Estimate plane by PCA of neighbors

Some Useful Software

- MeshLab (www.meshlab.net)
 - Viewing and processing meshes
- OpenMesh (www.openmesh.org)
 - Mesh processing
- CGAL (www.cgal.org)
 - Computational geometry