




Classical Geometry Processing

Prof. Angela Dai

Brief Recap

Machine Perception of Real-World Environments



ScanNet [Dai et al. 17]

We perceive and interact with a 3D world




ASIMO, Honda



Star Trek TNG (Phantasms)


How to represent 3D?



Discrete:
Meshes,
Point Samples



Parametric




Implicit:
Distance Fields

Poisson Surface Reconstruction

- State-of-the-art reconstruction
- Oriented points to implicit function
- Formulates reconstruction as a variational problem

$$\min_f \|\nabla f - V\|$$

where V is the vector field defined by the point samples



Poisson Surface Reconstruction


- No machine learning
- No data-driven priors
- How to fill in large holes?



[Kazhdan et al. '06, Kazhdan et al. '13]


Poisson Surface Reconstruction

- Great with dense point samples
- Can be sensitive to V
 - Normal estimates need to be reliable
 - If multiple sensor measurements are taken, alignment must be good
- No general or semantic priors




[Kazhdan et al. '06, Kazhdan et al. '13]

Registration and Alignment




Geometry Acquisition

- 3D Scanning and Motion Capture (IN2354)



3D Alignment

- Many applications, e.g., 3D scanning, SLAM




3D Alignment

- Many applications, e.g., 3D scanning, SLAM



3D Alignment

- Many applications




Protein Structure Alignment:

- (red) from humans
- (yellow) from fly *Drosophila melanogaster*.

3D Alignment (Registration)

- Input:
 - 2 shapes A and B with partial overlap
- Problem:
 - Register B to A by rigid transform, minimizing distance between A and B



How to measure distance?

- Measure of success for registration problem


$$\min_T \delta(A, T(B))$$

T : rigid transform to bring B to A

- Fundamental for geometric similarity, classification, general machine learning losses

How to evaluate 3D distance?



- What about common function norms, e.g., ℓ_2 ?
 - We don't have correspondences across 3D structures, shapes
- Should support partial matches



- Trade-off between support size and aggregated distance
- Distance for partial matches not a metric

Alignment Estimation

- Given shapes A and B
- Establish correspondences between A and B



- Find optimal transform that best aligns correspondences together, based on a distance measure

Transform Estimation

- Degrees of freedom
- Rigid motion has 6 degrees of freedom (3 rotation, 3 translation)
- Typically estimate with more correspondences -> overdetermined problem
- More general transforms -> more degrees of freedom, e.g., nonrigid deformations




Alignment Challenges

- Correspondence Estimation
 - Combinatorial search
- Transform Estimation
 - Transforms can be non-linear
- Difficult optimization -> look for good features, low-dimensional transforms


Rigid 3D Alignment

- Find 6DoF rigid transform that best aligns shapes, even if the shapes are different



Rigid 3D Alignment (Given Correspondences)

- Given correspondences $\{x_i\}, \{y_i\} \in \mathbb{R}^3$
- Find rigid transform \mathbf{R}, t that minimizes $\sum_{i=1}^N \|\mathbf{R}x_i + t - y_i\|_2^2$



Solved as *orthogonal Procrustes problem* in 1966



Rigid 3D Alignment (Given Correspondences)

$$\min_{\mathbf{R}, t} \sum_{i=1}^N \|\mathbf{R}x_i + t - y_i\|_2^2$$

- How to solve for \mathbf{R}, t ?
- Consider coordinate system centered at the mean of the x_i

$$\min_{\mathbf{R}, t} \underbrace{\sum_{i=1}^N \|t - y_i\|_2^2}_{\text{translation part}} - 2 \underbrace{\sum_{i=1}^N \langle \mathbf{R}x_i, y_i \rangle}_{\text{rotation part}}$$

Rigid 3D Alignment (Given Correspondences)

$$\min_{\mathbf{R}, t} \sum_{i=1}^N \|t - y_i\|_2^2 - 2 \sum_{i=1}^N \langle \mathbf{R} x_i, y_i \rangle$$

- Translation: $t = \frac{1}{N} \sum_{i=1}^N y_i$ (align centroids)
- Remove translation by mean-centering:


$$\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i \quad \bar{y} = \frac{1}{N} \sum_{i=1}^N y_i \quad X = [x_0 - \bar{x}, \dots, x_n - \bar{x}]^T \quad Y = [y_0 - \bar{y}, \dots, y_n - \bar{y}]^T$$

	N
3	X
	Y

- Compute SVD: $XY^T = UDV^T \leftarrow 3 \times 3$ matrix
- Define $S = \begin{cases} I, & \text{if } \det(U) \det(V) = 1 \\ \text{diag}(1, \dots, 1, -1) & \text{otherwise} \end{cases}$

$$R = USV^T$$

How to get correspondences?





How to get correspondences?

- Iterate between finding correspondences and solving for the best transform for those correspondences
 - Iterative Closest Points (ICP)
- Various methods to explicitly match features (handcrafted or learned), which can also be refined with ICP


Iterative Closest Points (Besl and McKay '92)

- Developed for aligning 3D shapes
- Nice analysis: *Efficient variants of the ICP algorithm* (Rusinkiewicz and Levoy 2001)




Iterative Closest Points

- How to find correspondences?
- Assume that closest points correspond
- Align the P_a points to their closest P_B neighbors; repeat
- Converges if starting positions are "close enough"




Iterative Closest Points




- Given a pair of shapes, A and B
- Iterate:
 - Find corresponding points P_A and P_B based on proximity
 - Find optimal transform \mathbf{R}, t minimizing $\operatorname{argmin}_{\mathbf{R}, t} \sum_i \|\mathbf{R}x_i + t - y_i\|_2^2$
 - Apply optimized \mathbf{R}, t

Iterative Closest Points



- Given a pair of shapes, A and B
- Iterate:
 - Find corresponding points P_A and P_B based on proximity
 - Find optimal transform \mathbf{R}, t minimizing $\operatorname{argmin}_{\mathbf{R}, t} \sum_i \|\mathbf{R}x_i + t - y_i\|_2^2$
 - Apply optimized \mathbf{R}, t

Iterative Closest Points




- Given a pair of shapes, A and B
- Iterate:
 - Find corresponding points P_A and P_B based on proximity
 - Find optimal transform \mathbf{R}, t minimizing $\operatorname{argmin}_{\mathbf{R}, t} \sum_i \|\mathbf{R}x_i + t - y_i\|_2^2$
 - Apply optimized \mathbf{R}, t

ICP: Runtime Analysis


- Each iteration:
 - Find closest points:
 - $O(N_B)$ per point
 - $O(N_B * N_A)$ total
 - Compute optimal alignment: $O(N_A)$
 - Update scene $O(N_A)$
- Speed up with fast or approximate nearest-neighbor data structures, e.g., *kd*-tree

ICP Analysis

- 
- Selection of points
 - Matching correspondences
 - Weighting correspondences
 - Rejecting outlier correspondences
 - Assigning error metric to the current transform
 - Minimizing error metric w.r.t transform

ICP Analysis

- How to select correspondences?




But: Uneven Sampling

Ideally, 1:1 correspondences

ICP Analysis

- How to select correspondences?




In practice: faster convergence than point-point ICP

ICP


- Solve for rigid transform between two surfaces
- Still used today
- Small basin of convergence
 - needs a good initialization
 - Hard to find correspondences
 - Need “sufficient” overlap

How to establish correspondences?



How to establish correspondence?



- When do two points on different shapes/scans represent the same feature?
- Are the surrounding regions similar?



Feature
descriptors
summarize
surrounding
regions


Classical Descriptors

- Curvature
- Differential features describe characteristics of surrounding surface
- Differential features can be noisy on meshes and real-world captured data



Classical Descriptors: Spin Images



- Create image associated with neighborhood of a feature point
- "Spin" image along point normal
- Collect contributions of each other point by their distance to tangent and distance to normal
- 2D spin image comparison



Spin Images [Johnson and Hebert '99]

Classical Descriptors: Point Feature Histograms


- For a point p find its k neighbors $\{q_i\}$
- Compute histogram from tuples of $\{(p, q_i)\}$ based on distances, normal, optionally curvature etc.



Classical Descriptors

- Often fast to compute
- Designed to be used with Euclidean distances, can use accelerated search structures for fast matching
- In practice: often have RGB images
 - Use classical RGB features (e.g., SIFT, SURF, etc.)

Non-rigid Deformations





Rigid vs Deformable Registration

- Rigid: uniformly applied rigid transformation
 - Rotation, translation
 - Maintains same shape and size
- Deformable: allows non-uniform mapping between surfaces
 - Applications: to correct for small distortions in measurement
 - Modeling non-rigid shape movement over time


Non-Rigid Shape Matching

- Consider near-isometric cases
- Find correspondences that preserve intrinsic (geodesic) distances on the shapes



Measuring intrinsic shape similarity



- Gromov-Hausdorff distance
- Hausdorff distance: maximin
 - Maximum of all minimum distances between two sets of points



- Gromov-Hausdorff: infimum of all Hausdorff distances over mappings or correspondences
 - Over all correspondences -> difficult to compute!

Near isometries preserve local structure


- Optimal alignment can be defined, but difficult to compute
- Define descriptors of local regions -> establish good mappings




How large
should a local
region be?

Scale for local features?

- Given a point on a shape, find other points with similar neighborhoods



smaller scale



larger scale


- How to meaningfully compare neighborhoods at different scales?

Heat kernel signature

- Spectral shape analysis
- Heat kernel $k_t(x, y)$: amount of heat transferred from x to y in time t

$$f(x, t) = \int_M k_t(x, y) f(y) dy$$


- Invariant under isometric deformations
- Multi-scale description



- In practice: can be difficult to apply to real scenarios (noisy, partial data)

As-Rigid-As-Possible (ARAP)


- Idea: maximize local rigidity



As-Rigid-As-Possible (ARAP)


- Idea: maximize local rigidity
- Ask all star edges to transform rigidly by some rotation R
 - Shape of cell is preserved

$$\min \sum_{j \in N(i)} \|(v'_i - v'_j) - R_i(v_i - v_j)\|^2$$



As-Rigid-As-Possible (ARAP)

- If v, v' are known, then R_i is uniquely defined



- Build covariance matrix $S = VV'^T$
- SVD: $S = U\Sigma W^T, R_i = UW^T$

ARAP Deformation Energy

$$\min \sum_{i=1}^n \sum_{j \in N(i)} \|(v'_i - v'_j) - R_i(v_i - v_j)\|^2$$


- Variables: v', R
- Can iteratively optimize alternating solving for v', R
 - Solve for R_i with fixed v, v_i
 - With fixed R_i , minimize energy by finding $v': Lv' = b$

ARAP

- Good at preserving edge lengths
- Fast on small meshes
- No understanding of volume
- Slow propagation on large meshes



Scene Flow

- 3D motion field of points in the world



Scene Flow

- Between two 3D observations P_t, P_{t+1}
- Flow is vector field Δ_t ; for each point in P_t defining the motion to P_{t+1}
- $P_{t+1}(x, y, z) = P_t(x + \Delta x, y + \Delta y, z + \Delta z)$



Additional References

- Efficient variants of the ICP algorithm [Rusinkiewicz et al. '01]
 - http://www.pcl-users.org/file/n4037867/Rusinkiewicz_Efficient_Variants_of_ICP.pdf
- As-Rigid-As-Possible Surface Modeling [Sorkine et al '07]
 - <https://igl.ethz.ch/projects/ARAP/>